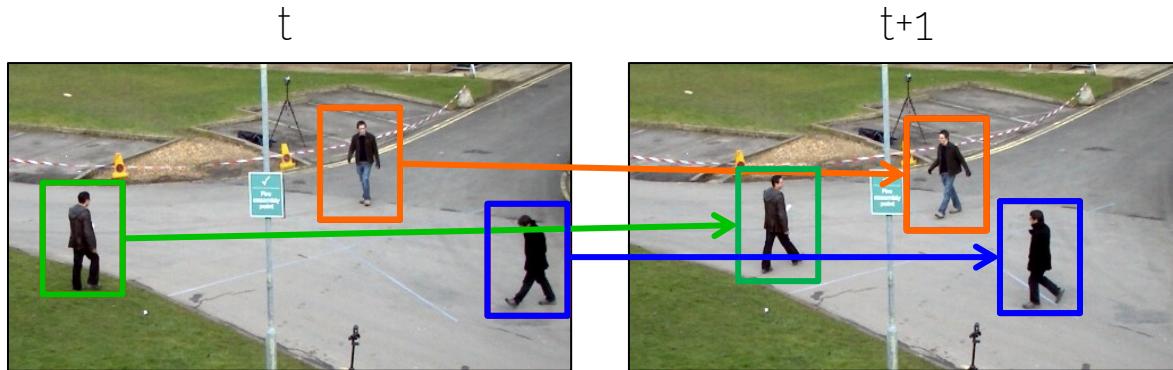


Object tracking

Problem statement

- Given a video, find out which parts of the image depict the same object in different frames
- Often we use detectors as starting points



Why do we need tracking?

- To model objects when detection fails:
 - Occlusions
 - Viewpoint/pose/blur/illumination variations (in a few frames of a sequence)
 - Background clutter
- To reason about the dynamic world, e.g., trajectory prediction (is the person going to cross the street?)

Tracking is....

- Similarity measurement
- Correlation
- Correspondence
 - Story time: „A young graduate student asked Takeo Kanade what are the three most important problems in computer vision. Kanade replied: “Correspondence, correspondence, correspondence!”
- Matching/retrieval
- Data association

Tracking is also....

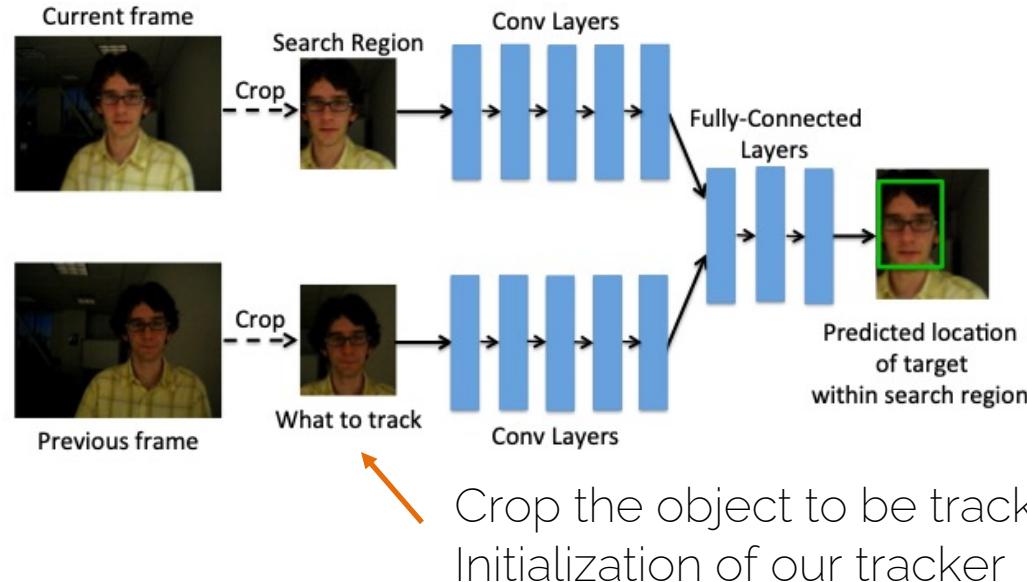
- Learning to model our target's
- Appearance: we need to know how the target looks like
 - Single object tracking
 - Re-identification
- Motion: to make predictions of where the targets goes
 - Trajectory prediction (lecture 6)

Single Target Tracking

- STT (1) as a matching/correspondence problem:
 - GOTURN: no online appearance modeling
- STT (2) as an appearance learning problem:
 - MDNet: quick online finetuning of the network
- STT (3) as a (temporal) prediction problem:
 - ROLO = CNN + LSTM

Single Target Tracking 1

- Input: what to track?

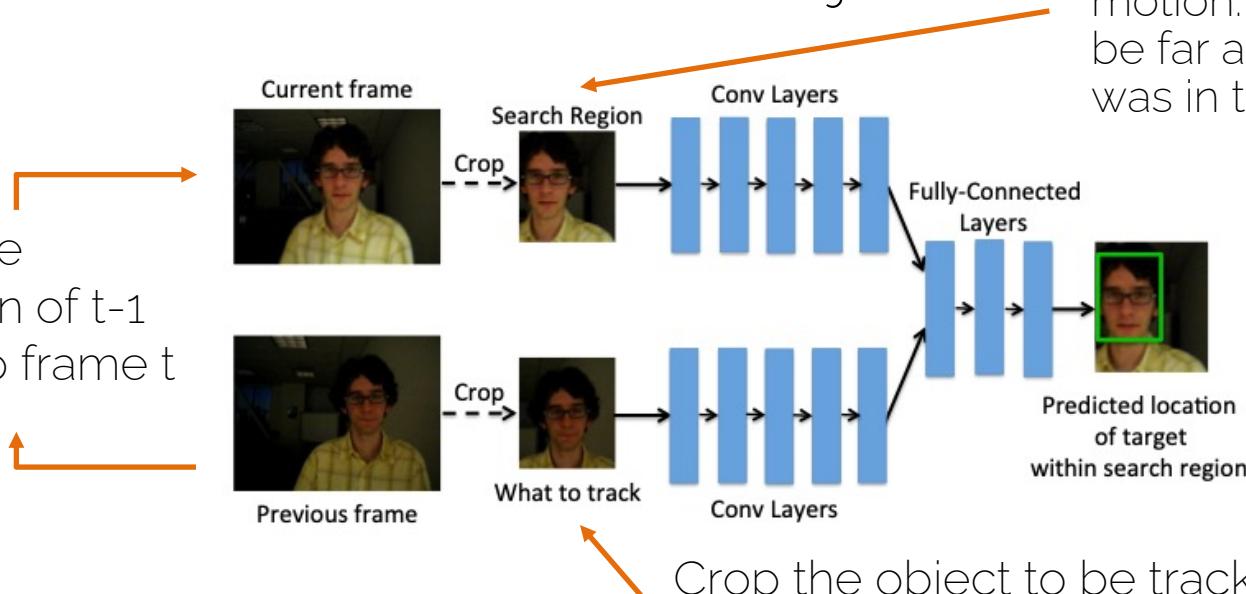


D. Held, S. Thrun, S. Savarese. "Learning to Track at 100 FPS with Deep Regression Networks". ECCV 2016.

Single Target Tracking 1

- Where do I search for the object?

Use the position of $t-1$ to crop frame t



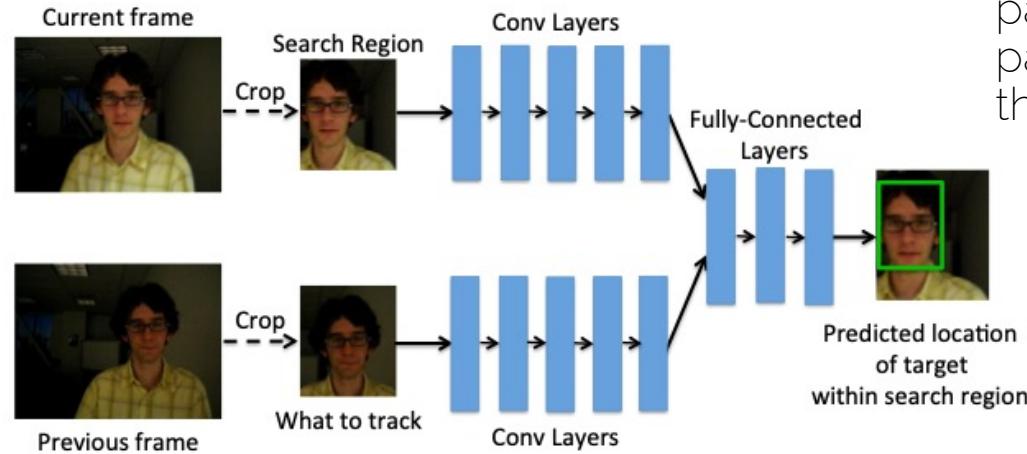
Assume smooth "slow" motion. The object cannot be far away from where it was in the previous frame.

Crop the object to be tracked –
Initialization of our tracker

D. Held, S. Thrun, S. Savarese. "Learning to Track at 100 FPS with Deep Regression Networks". ECCV 2016.

Single Target Tracking 1

- Architecture: conv + concatenate + FC



Check the original paper for the exact parameterization of the output

D. Held, S. Thrun, S. Savarese. "Learning to Track at 100 FPS with Deep Regression Networks". ECCV 2016.

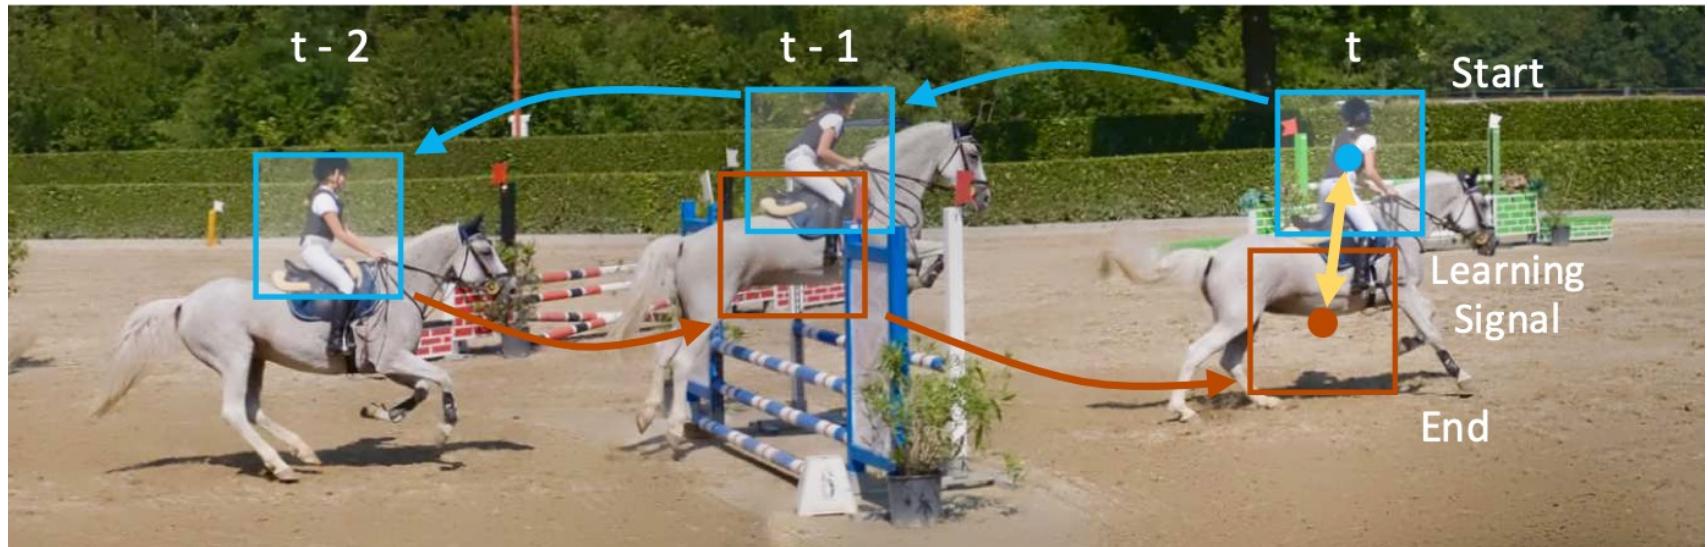
Single Target Tracking 1

- PROS of GOTURN:
 - No online training required.
 - Tracking is done by comparison, so we do not need to retrain or finetune our model for every new object.
 - Close to the template matching approach that we saw in the first lectures for object detection
 - This makes it very fast!
- CONS:
 - We have a motion assumption. If the object moves fast and goes out of our search window, we cannot recover.

D. Held, S. Thrun, S. Savarese. "Learning to Track at 100 FPS with Deep Regression Networks". ECCV 2016.

SOT 1.2 - Unsupervised

- Forward cycle and backward cycle should be consistent!



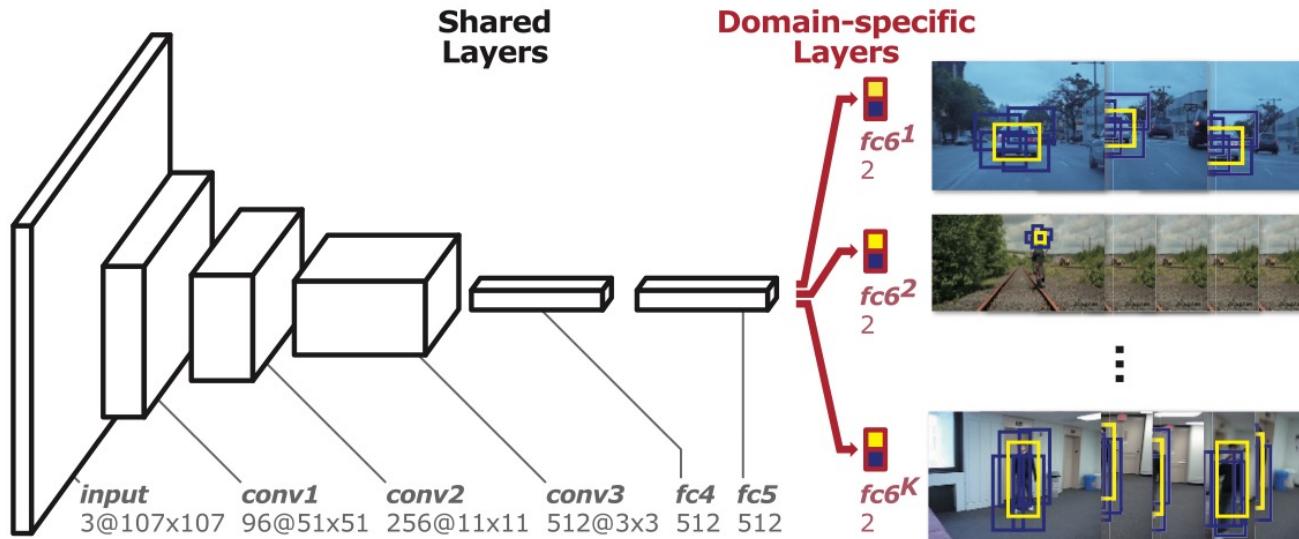
X. Wang, A. Jabri, A. Efros. "Learning correspondence from the cycle-consistency of time". CVPR 2019

Single Target Tracking 2

- Online appearance model learning entails training your CNN at test time.
 - Slow: not suitable for real-time applications
 - Solution: train as few layers as possible

Single Target Tracking 2

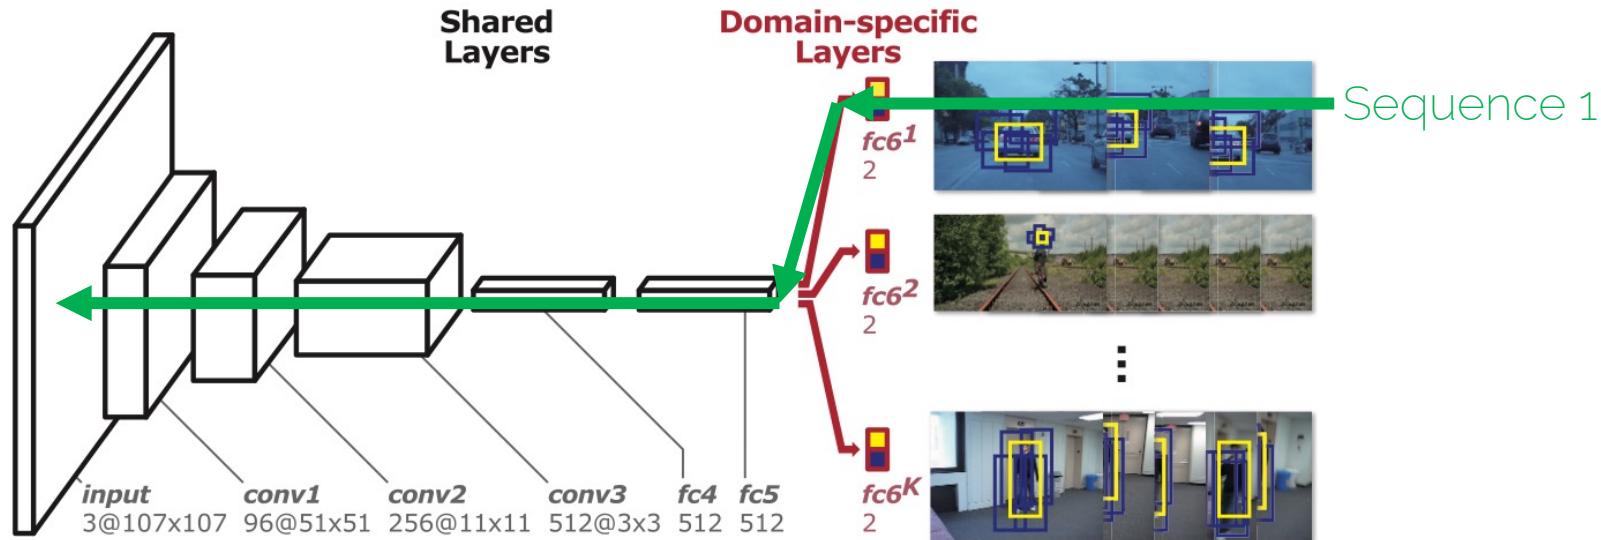
- Shared layers + scene-specific layers



H. Nam and B. Han, „Learning Multi-Domain Convolutional Neural Networks for Visual Tracking“. CVPR 2016

Single Target Tracking 2

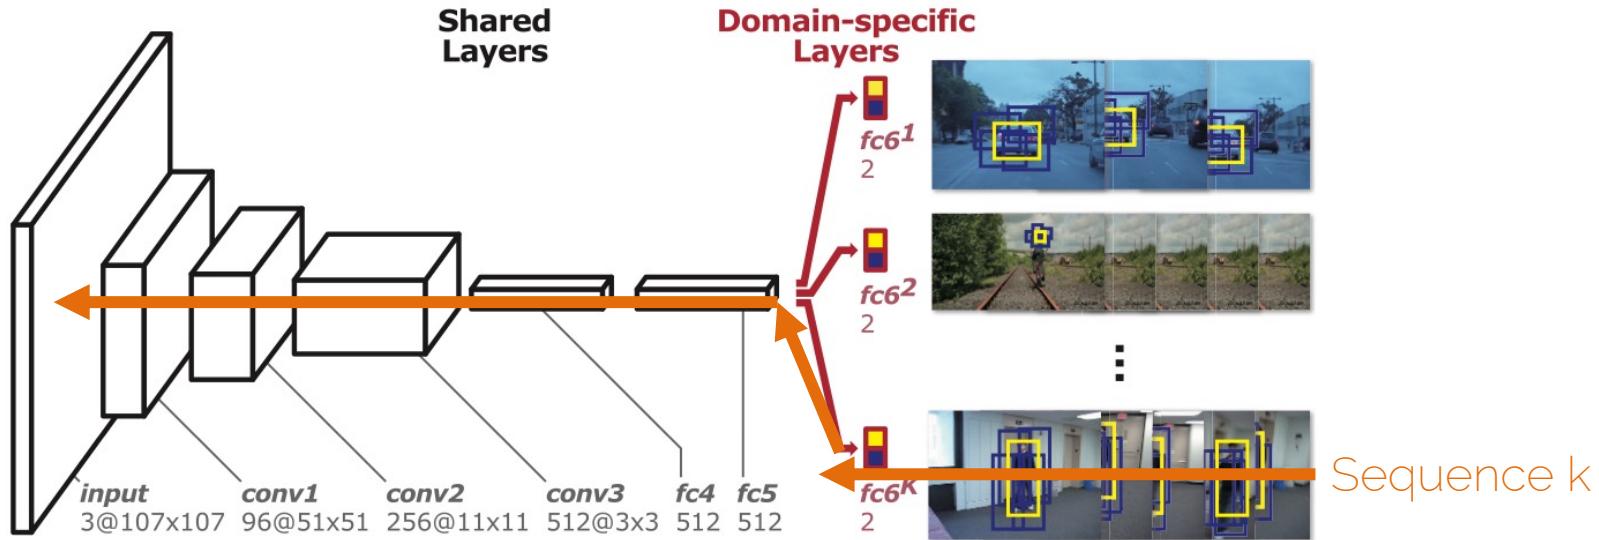
- Backpropagation is independent per sequence



H. Nam and B. Han, „Learning Multi-Domain Convolutional Neural Networks for Visual Tracking“. CVPR 2016

Single Target Tracking 2

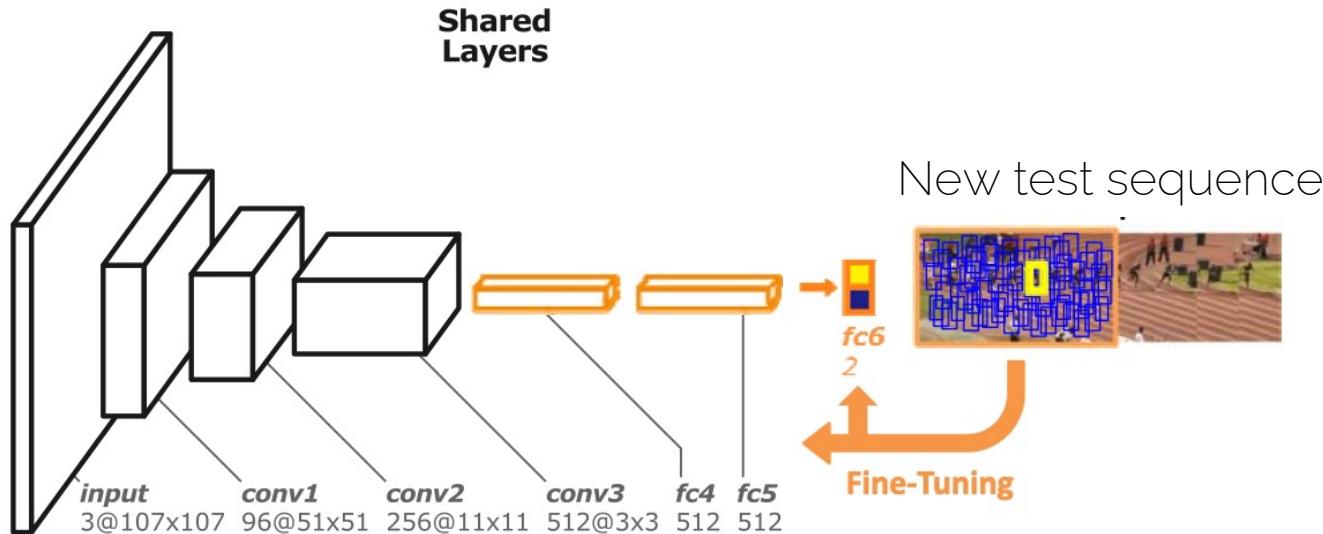
- Backpropagation is independent per sequence



H. Nam and B. Han, „Learning Multi-Domain Convolutional Neural Networks for Visual Tracking“. CVPR 2016

Single Target Tracking 2

- At test time, we need to train fc6 (up to fc4 if wanted).

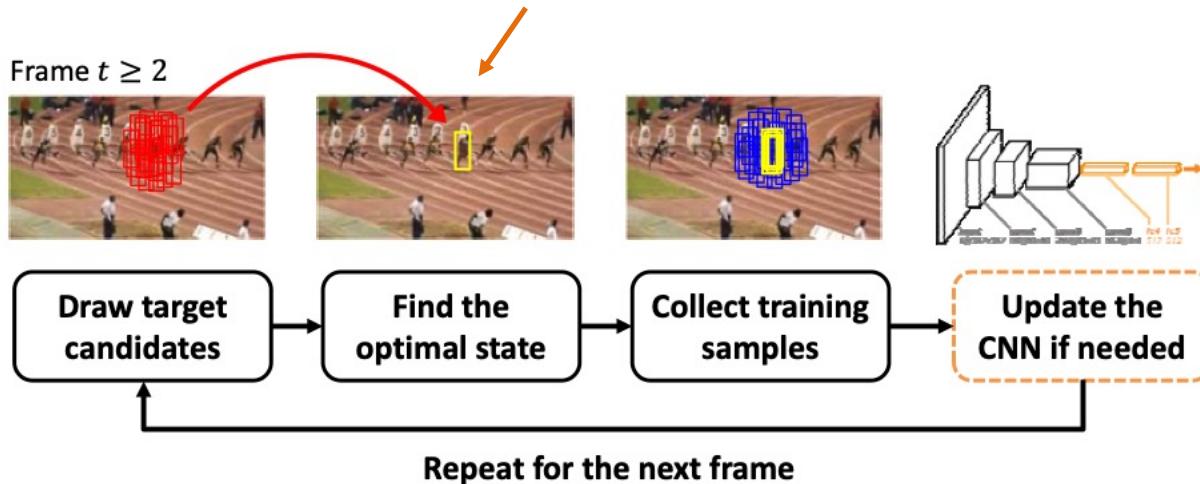


H. Nam and B. Han, „Learning Multi-Domain Convolutional Neural Networks for Visual Tracking“. CVPR 2016

Single Target Tracking 2

- Online tracking

R-CNN type of regression



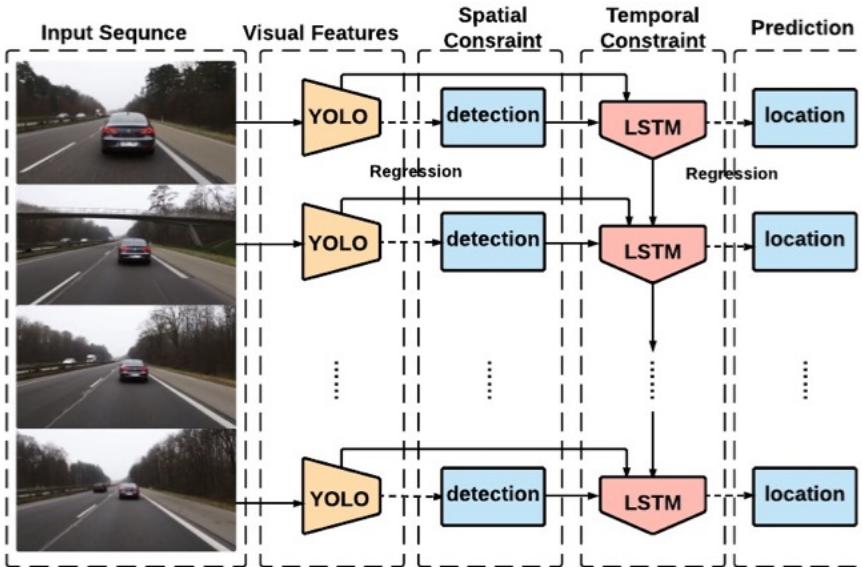
H. Nam and B. Han. „Learning Multi-Domain Convolutional Neural Networks for Visual Tracking“. CVPR 2016

Single Target Tracking 2

- PROS of MDNet:
 - No previous location assumption, the object can move anywhere in the image
 - Fine-tuning step is comparatively cheap
 - Winner of the VOT Challenge 2015 (<http://www.votchallenge.net>)
- CONS:
 - Not as fast as GOTURN

Single Target Tracking 3

- CNN for appearance + LSTM for motion

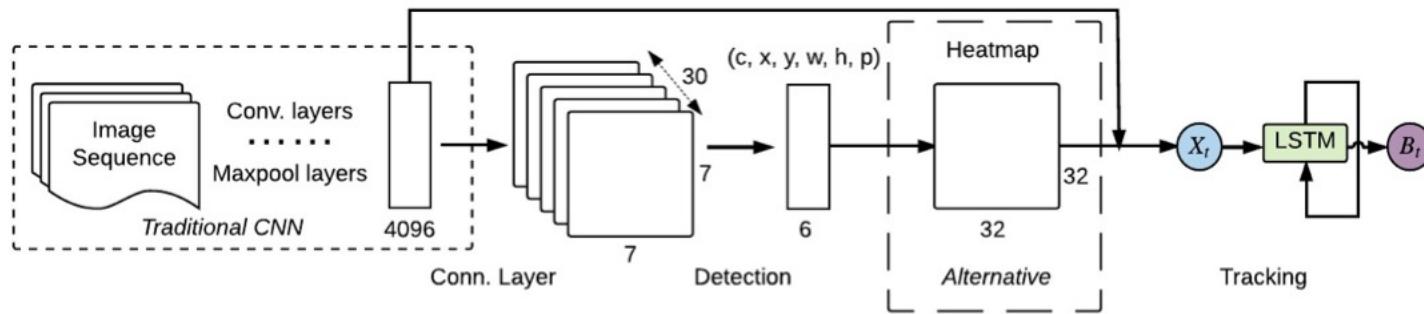


Recurrent YOLO

ROLO

Single Target Tracking 3

- LSTM receives the heatmap for the object's position and the 4096 descriptor of the image

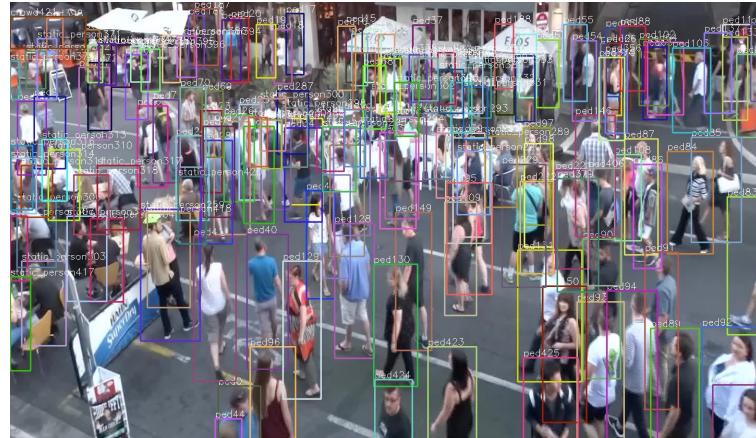
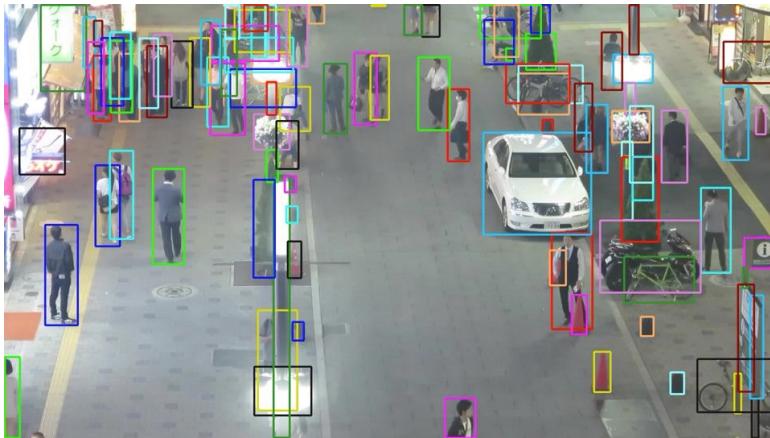


G. Ning, Z. Zhang, C. Huang, Z. He. „Spatially Supervised Recurrent Convolutional Neural Networks for Visual Object Tracking“. arXiv:1607.05781. 2016

Multiple object tracking

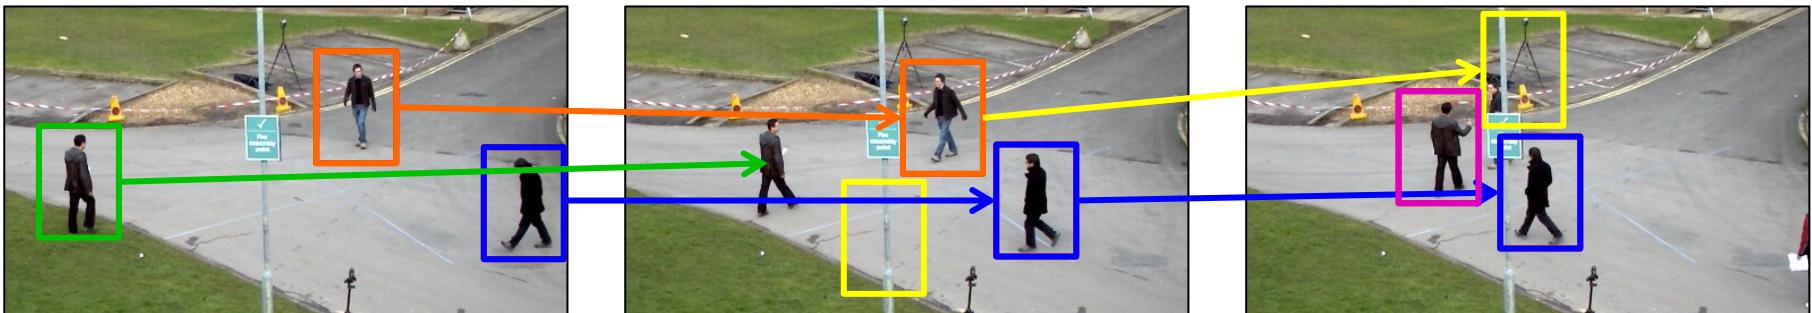
Different challenges

- Multiple objects of the same type
- Heavy occlusions
- Appearance is often very similar



Tracking-by-detection

- We will focus on algorithms where a set of detections is provided
 - Remember detections are not perfect!

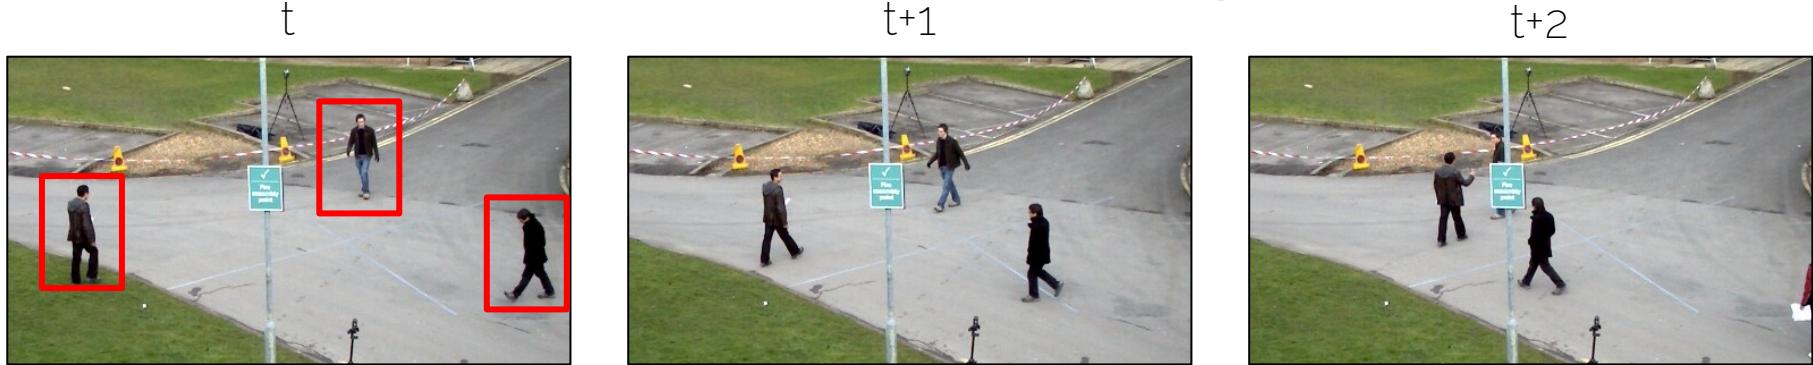


Find detections that match and form a trajectory

Online vs offline tracking

- Online tracking
 - Processes two frames at a time
 - For real-time applications
 - Prone to drifting → hard to recover from errors or occlusions
- Offline tracking
 - Processes a batch of frames
 - Good to recover from occlusions (short ones as we will see)
 - Not suitable for real-time applications
 - Suitable for video analysis

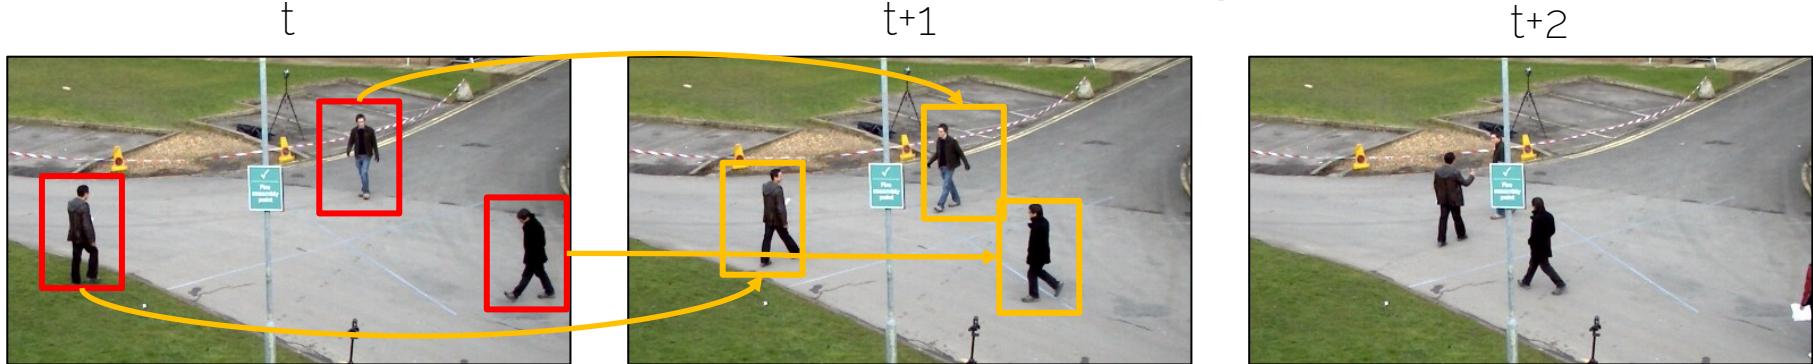
Online tracking



- 1. Track initialization (e.g. using a detector)



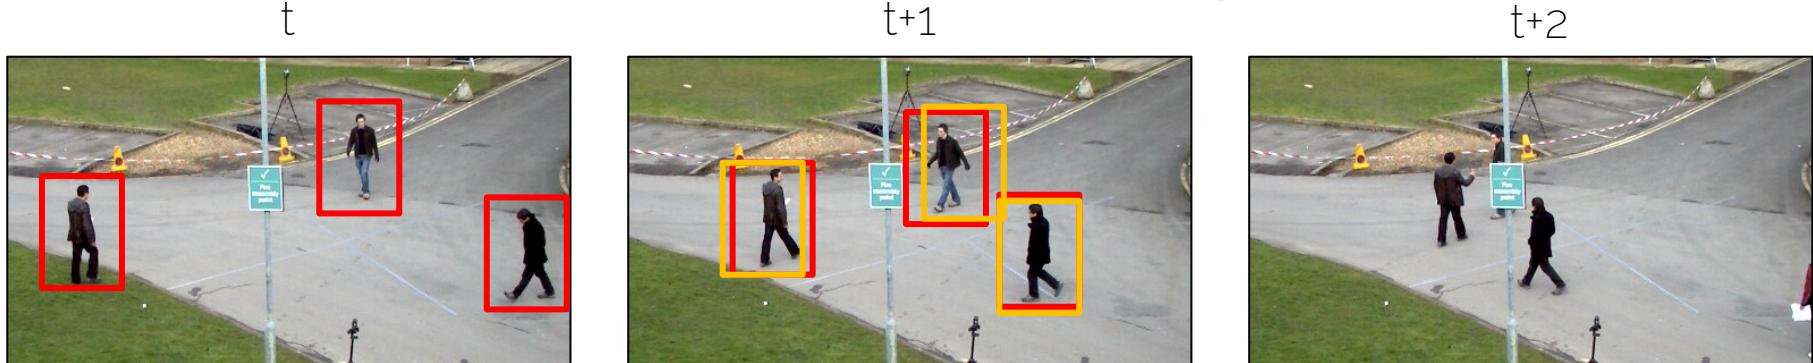
Online tracking



- 1. Track initialization (e.g. using a detector)

- 2. Prediction of the next position (motion model) 

Online tracking



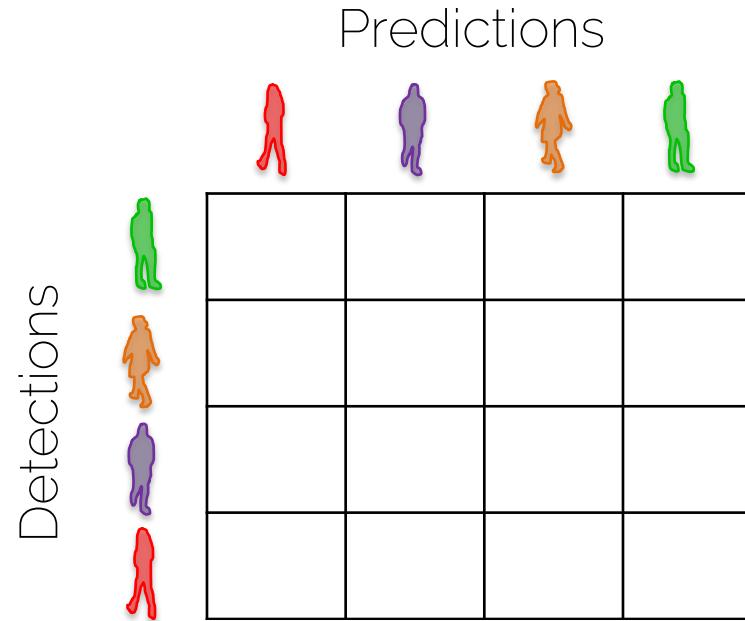
- 1. Track initialization (e.g. using a detector)
- 2. Prediction of the next position (motion model)
- 3. Matching predictions with detections (appearance model)

Online tracking

- 2. Prediction of the next position (motion model)
 - Classic: Kalman filter
 - Nowadays: Recurrent architecture
 - For now: we will assume a constant velocity model (spoiler alter: it works really well at high framerates and without occlusions!)

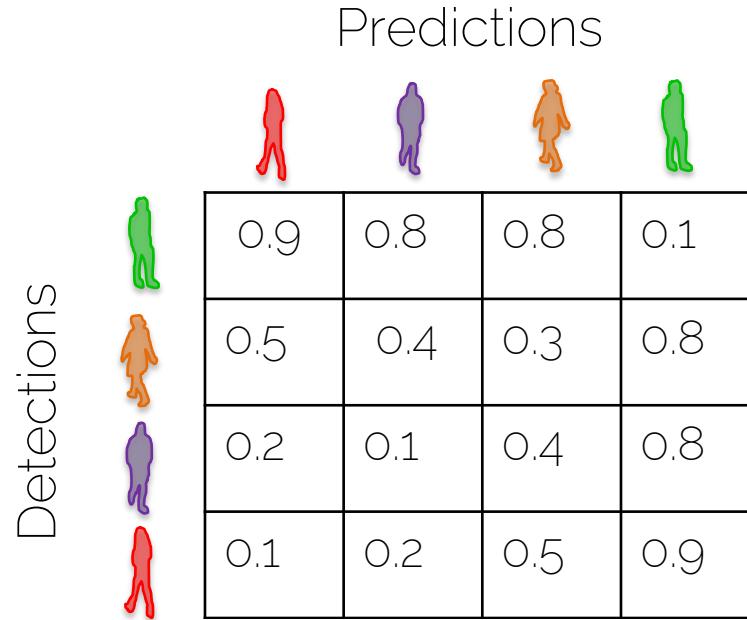
Online tracking

- 3. Matching predictions with detections



Online tracking

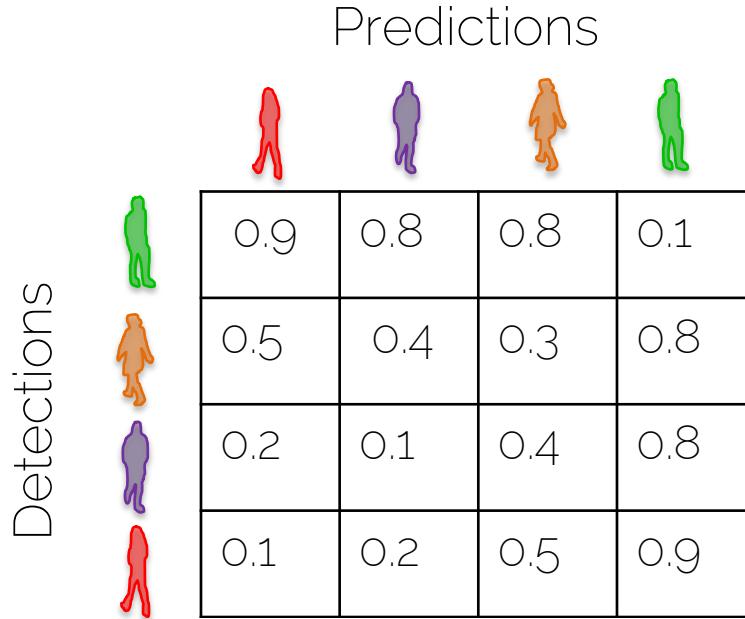
- Bipartite matching
 - Define distances between boxes (e.g., IoU, pixel distance, 3D distance)



Online tracking

- Bipartite matching

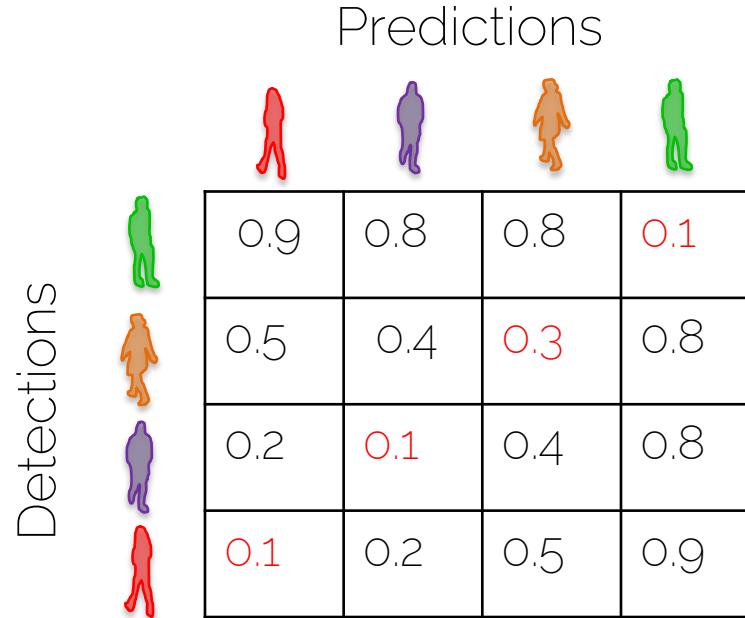
- Define distances between boxes (e.g., IoU, pixel distance, 3D distance)
 - Solve the unique matching with e.g., the Hungarian algorithm*



*Demo: <http://www.hungarianalgorithm.com/solve.php>

Online tracking

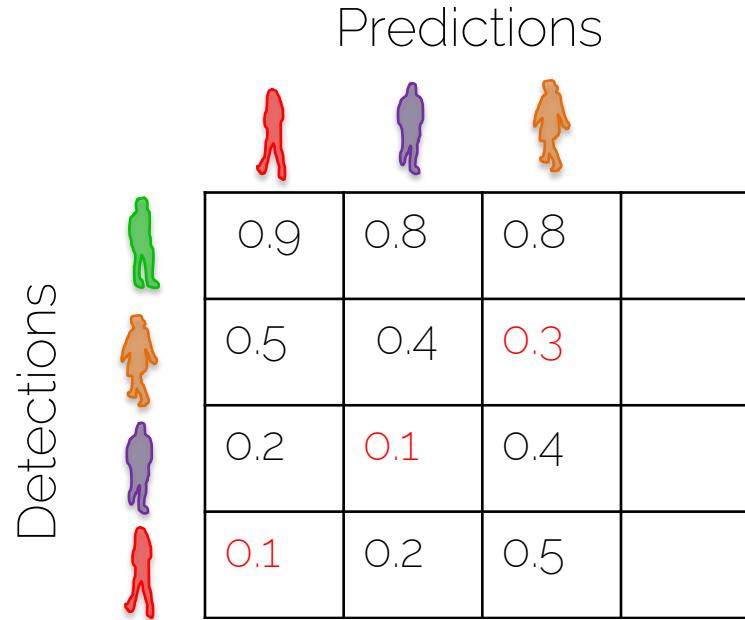
- Bipartite matching
 - Define distances between boxes (e.g., IoU, pixel distance, 3D distance)
 - Solve the unique matching with e.g., the Hungarian algorithm*
 - Solutions are the unique assignments that minimize the total cost



*Demo: <http://www.hungarianalgorithm.com/solve.php>

Online tracking

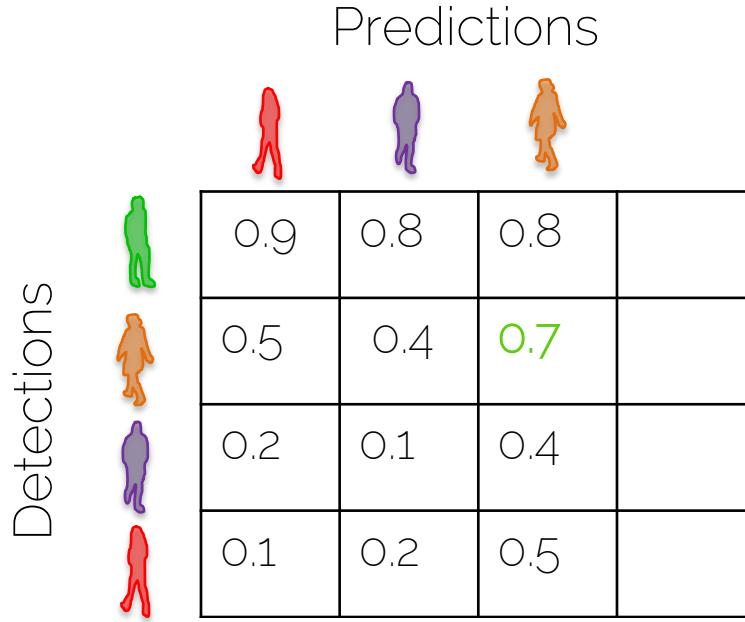
- Bipartite matching
 - What happens if we are missing a prediction?



*Demo: <http://www.hungarianalgorithm.com/solve.php>

Online tracking

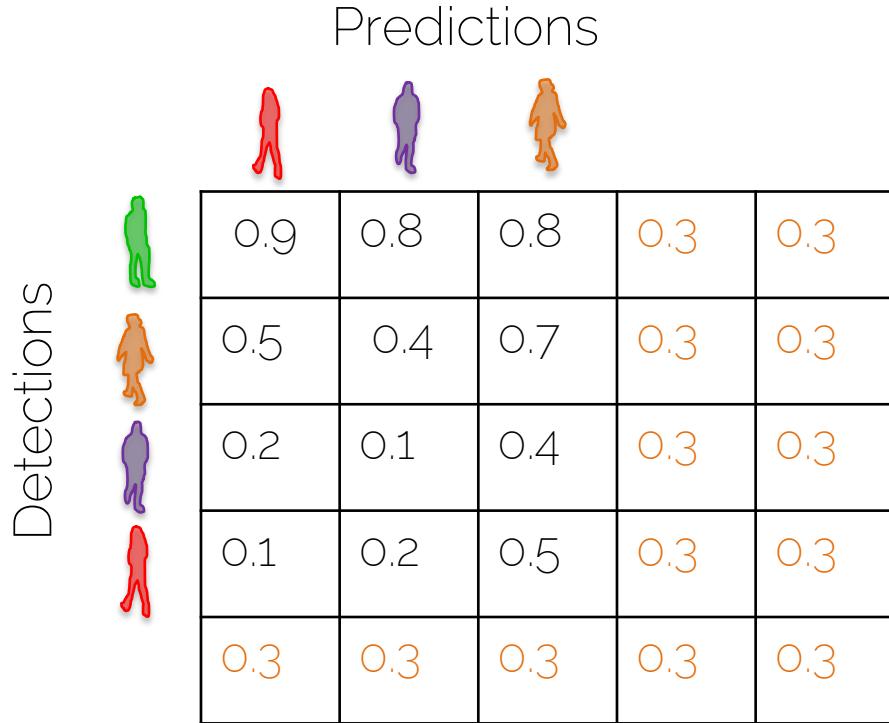
- Bipartite matching
 - What happens if we are missing a prediction?
 - What happens if no prediction is suitable for the match?



*Demo: <http://www.hungarianalgorithm.com/solve.php>

Online tracking

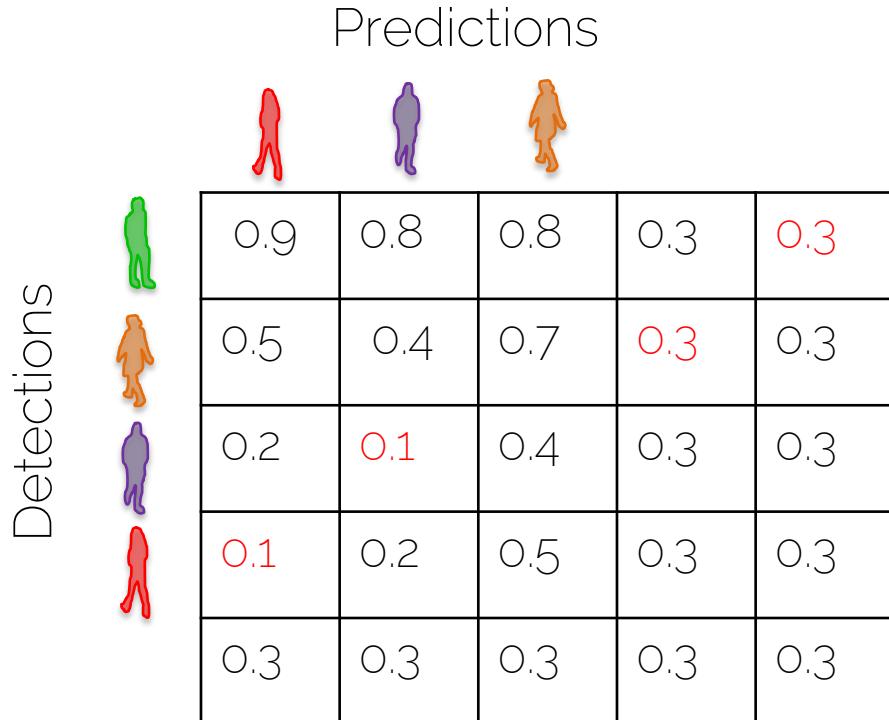
- Bipartite matching
 - What happens if we are missing a prediction?
 - What happens if no prediction is suitable for the match?
 - Introducing extra nodes with a threshold cost



*Demo: <http://www.hungarianalgorithm.com/solve.php>

Online tracking

- Bipartite matching
 - What happens if we are missing a prediction?
 - What happens if no prediction is suitable for the match?
 - Introducing extra nodes with a threshold cost
 - Apply Hungarian
 - Result: two detections have no matched prediction



*Demo: <http://www.hungarianalgorithm.com/solve.php>

Online tracking: what is DL's role?

- 1. Track initialization (e.g. using a detector)
 - Deep Learning has provided us with better detectors ✓
- 2. Prediction of the next position (motion model)
 - Trajectory prediction will be covered in Lecture 6 Adding temporal complexity
- 3. Matching predictions with detections (appearance model)
 - Improving appearance models → Re-Identification (in a few slides) Adding feature complexity
 - Matching still happens separately from learning, we will see how to couple both steps in the next lecture → MOT as a graph problem
Adding computational complexity

Tracktor: striving for simplicity

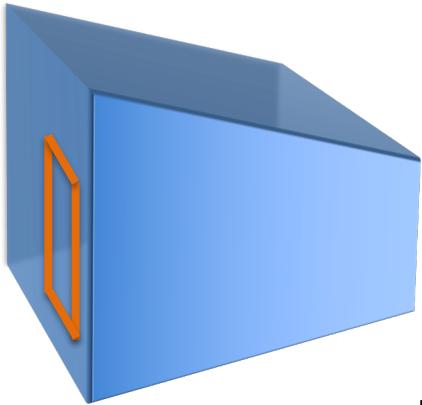
Recall two step-detectors

Input image

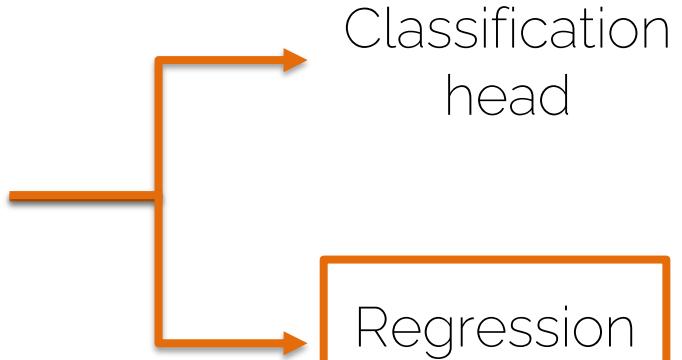


Region
proposal

Convolutions



Feature
representation



Classification
head

Regression
head

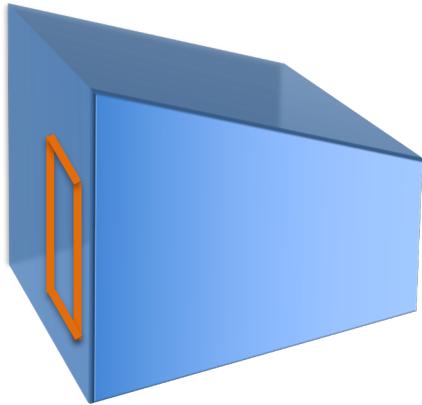
Recall two step-detectors

Input image

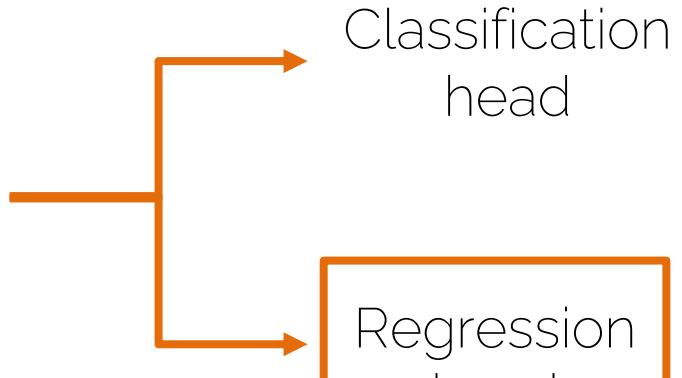


Region
proposal

Convolutions



Feature
representation



Classification
head

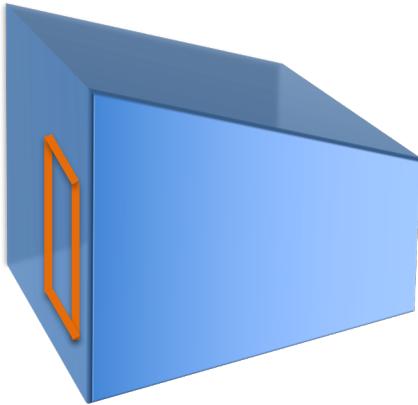
Regression
head

Recall two step-detectors

Input image

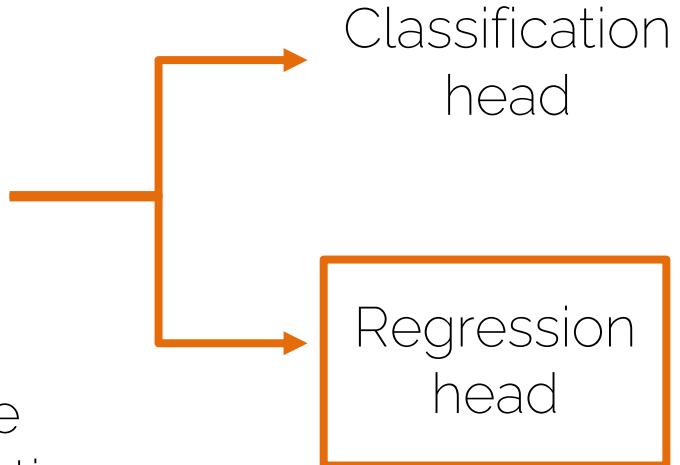


Convolutions



Regressed
bounding
box

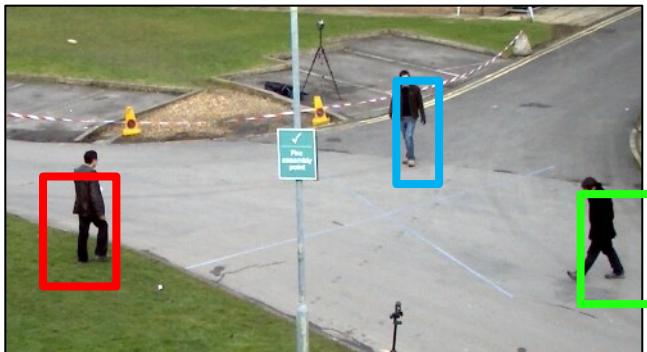
Feature
representation



Making a detector into a tracktor

- Tracktor: a method trained as a detector but with tracking capabilities

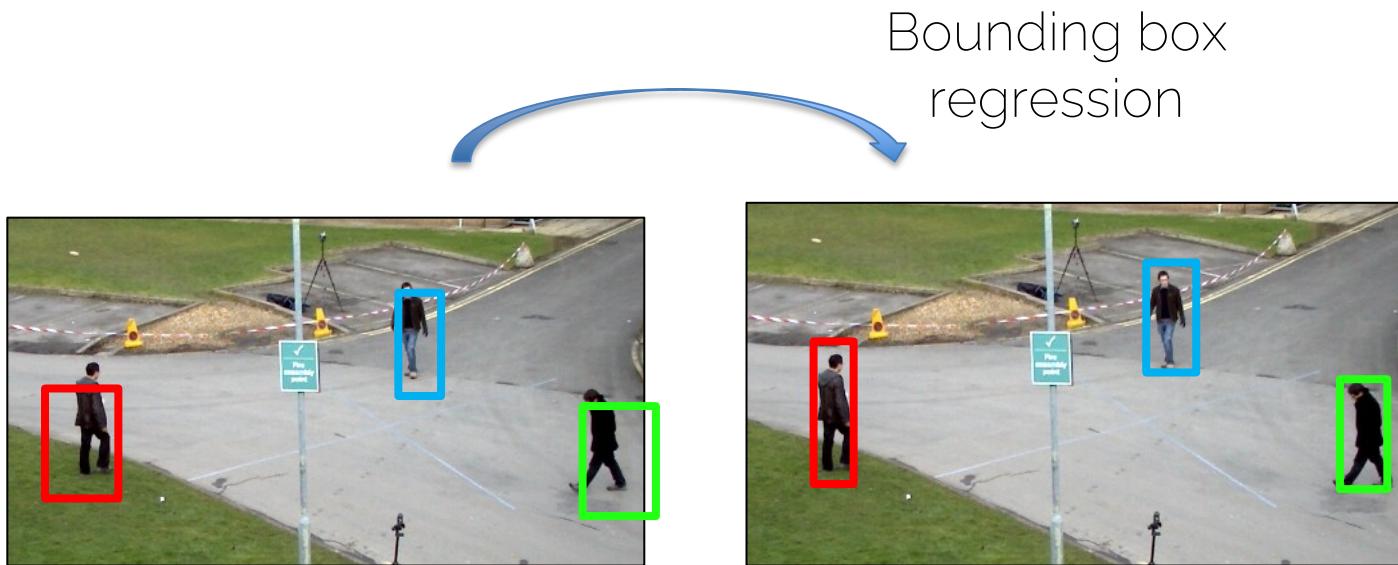
Making a detector into a tracktor



Frame $t+1$

Use detections of
frame t as proposals

Making a detector into a tracktor



Where did the detection with ID 1 go in the next frame? Tracking!

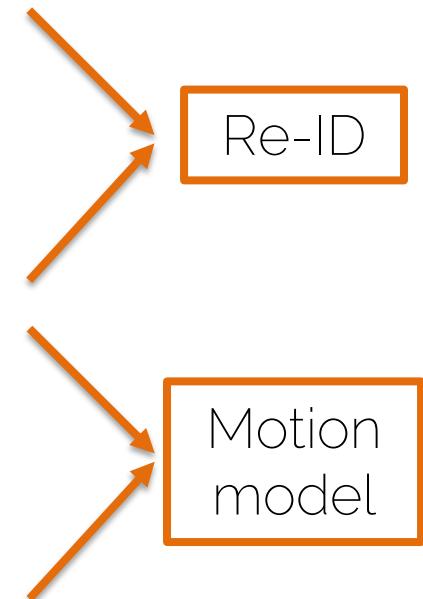
P. Bergmann, T. Meinhardt and L. Leal-Taixé. ICCV, 2019

Pros and cons

- PRO We can reuse an extremely well-trained regressor
 - We get well-positioned bounding boxes
- PRO We can train our model on still images → easier annotation!
- PRO Tracktor is online

Pros and cons

- CON There is no notion of "identity" in the model
 - Confusion in crowded spaces
- CON As any online tracker, the track is killed if the target becomes occluded
 - Need to close small gaps and occlusions
- CON The regressor only shifts the box by a small quantity
 - Large camera motions
 - Large displacements due to low framerate



Back to....

Modeling appearance



Re-ID

Modeling motion



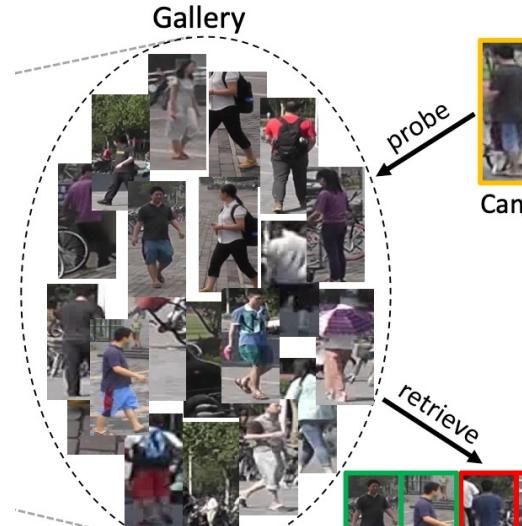
Motion
model

Re-ID

Problem statement

- Viewing tracking as a retrieval problem

Images of different identities



Detected person

...
Retrieve
the top
matches

How to measure (learn) the distance between two images?

Can we use classification?

- We can do image classification to get attributes

A



Classification: person, face, female, brunette

B



Classification: person, face, female, blonde

(Deep) similarity learning

- What if I just want to answer the question

A



Is it the same person?

B



(Deep) similarity learning

- Another type of problems:

A



B



- Similarity Learning: learn a function that measure how similar two objects are.
- Deep Metric Learning: Learning a distance function over objects.

Similarity Learning: when and why?

- Application: unlocking your iPhone with your face

Training



Similarity Learning: when and why?

- Application: unlocking your iPhone with your face

A



YES

B



NO

Testing



Can be solved as a classification problem

Similarity Learning: when and why?

- Application: face recognition system so students can enter the exam room without the need for ID check

Person 1



Training

Person 2



Person 3



Similarity Learning: when and why?

- Application: face recognition system so students can enter the exam room without the need for ID check

What is the problem
with this approach?

Scalability – we need to retrain our model every time a new student is registered to the course

Similarity Learning: when and why?

- Application: face recognition system so students can enter the exam room without the need for ID check

Can we train one
model and use it every
year?

Similarity Learning: when and why?

- Learn a similarity function

A



B



Low similarity
score (large
distance)

A



B



High similarity
score (small
distance)

Similarity Learning: when and why?

- Learn a similarity function: testing

A



$$d(A, B) > \tau$$

Not the same
person

B



Similarity Learning: when and why?

- Learn a similarity function

Same person

$$d(A, B) < \tau$$

A

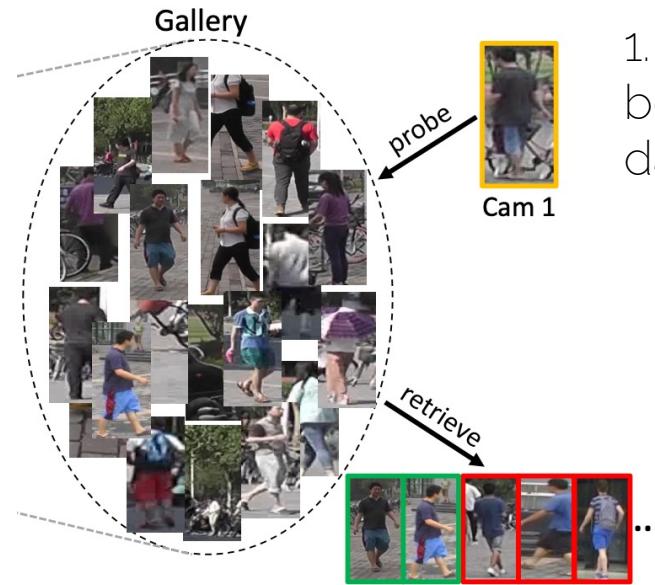


B



Similarity Learning: when and why?

- How about retrieval and re-ID?



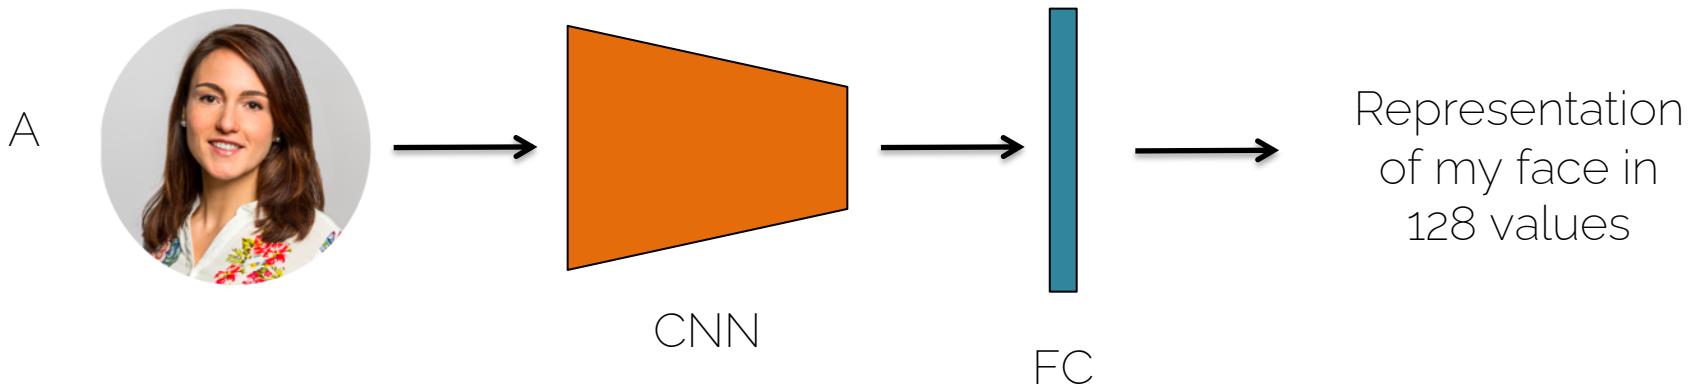
1. Compute distance between query image and database images
2. Use k-nearest neighbors, i.e., retrieve top images based on distance

Similarity learning

- How do we train a network to learn similarity?

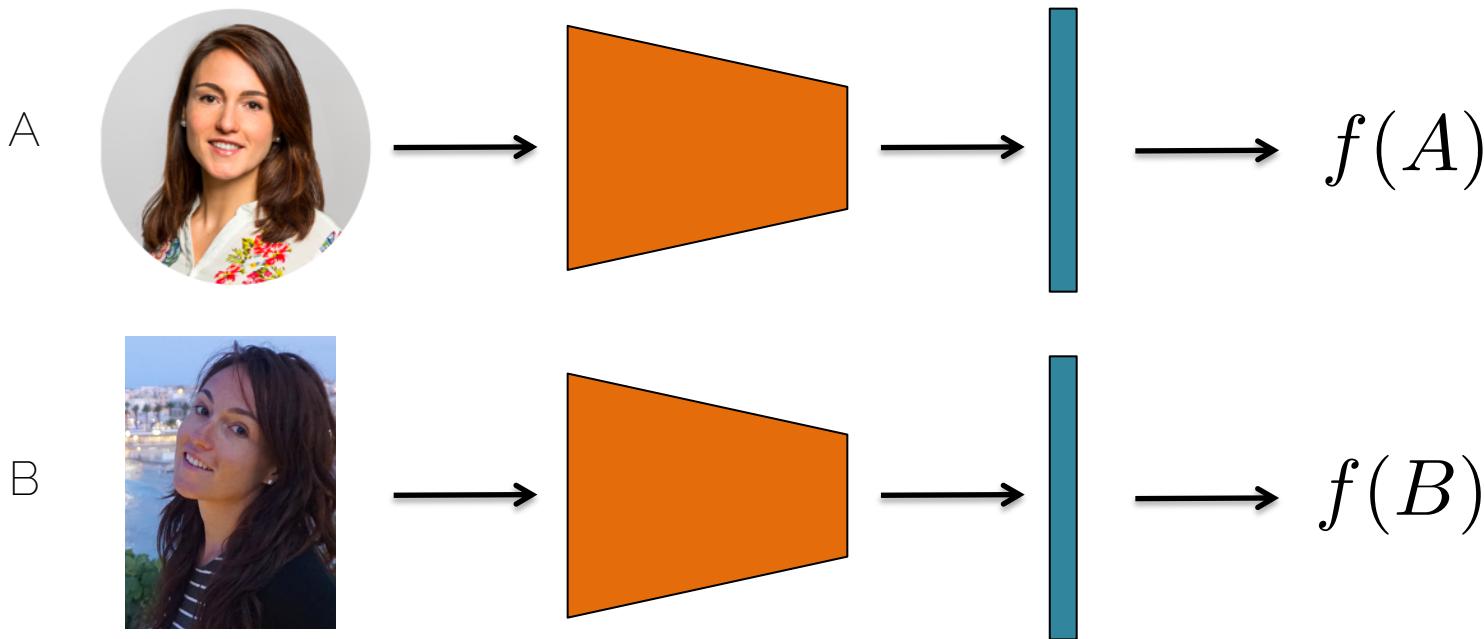
Similarity learning

- How do we train a network to learn similarity?



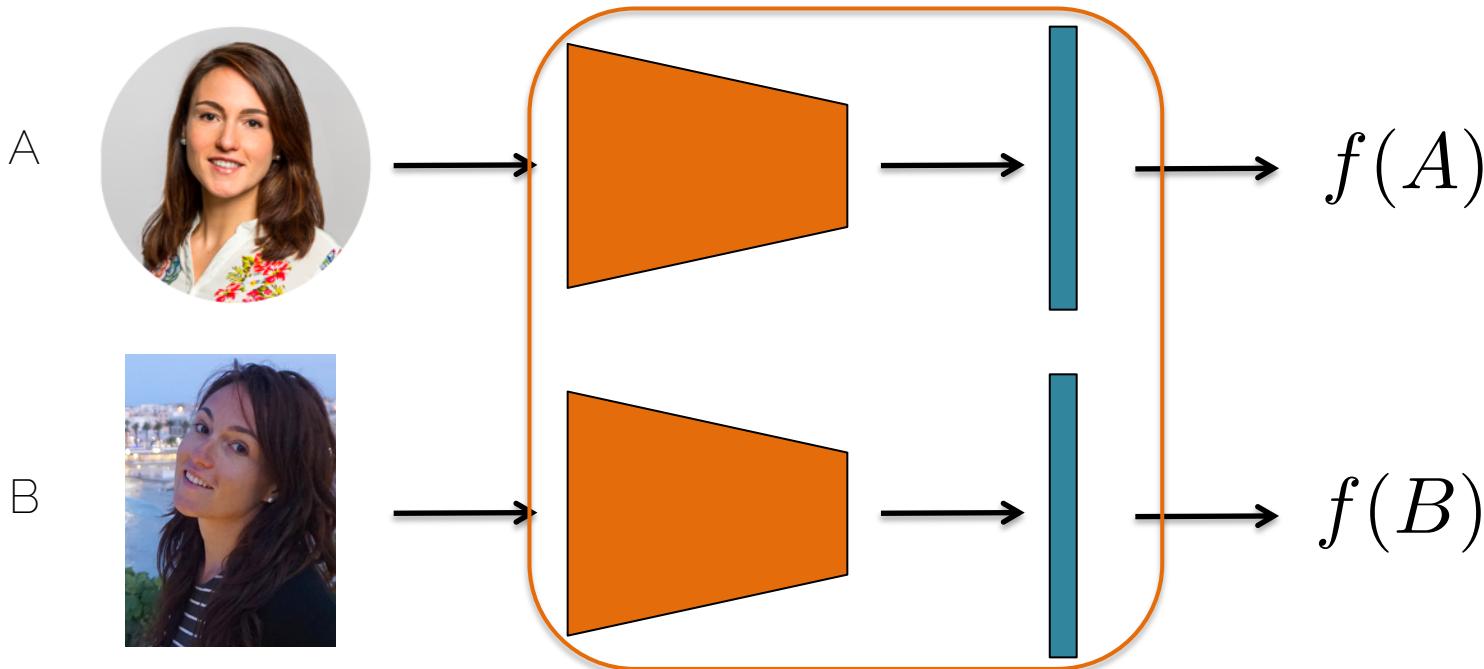
Similarity learning

- How do we train a network to learn similarity?



Similarity learning

- Siamese network = shared weights



Taigman et al. „DeepFace: closing the gap to human level performance“. CVPR 2014

Similarity learning

- Siamese network = shared weights
- We use the same network to obtain an encoding of the images, $f(A)$ and $f(B)$ respectively.
- To be done: compare the encodings

Similarity learning: losses

- Distance function $d(A, B) = ||f(A) - f(B)||^2$
- Training: learn the parameter such that
 - If A and B depict the same person, $d(A, B)$ is small
 - If A and B depict a different person, $d(A, B)$ is large

Similarity learning: losses

- Loss function for a positive pair:
 - If A and B depict the same person, $d(A, B)$ is small

$$\mathcal{L}(A, B) = ||f(A) - f(B)||^2$$

Similarity learning: losses

- Loss function for a negative pair:
 - If A and B depict a different person, $d(A, B)$ is large
 - Better use a Hinge loss:
$$\mathcal{L}(A, B) = \max(0, m^2 - ||f(A) - f(B)||^2)$$



margin

If two elements are already far away, do not spend energy in pulling them even further away

Similarity learning: losses

- Contrastive loss:

$$\mathcal{L}(A, B) = y^* \lVert f(A) - f(B) \rVert^2 + (1 - y^*) \max(0, m^2 - \lVert f(A) - f(B) \rVert^2)$$



Positive pair,
reduce the distance
between the
elements



Negative pair,
brings the elements
further apart up to a
margin

Chopra, Hadsell and LeCun. "Learning a similarity metric discriminatively, with application to Face verification", CVPR 2005.

Triplet loss

- Triplet loss allows us to learn a ranking



Anchor (A)



Positive (P)



Negative (N)

We want: $||f(A) - f(P)||^2 < ||f(A) - f(N)||^2$

Schroff et al „FaceNet: a unified embedding for face recognition and clustering“. CVPR 2015

Triplet loss

- Triplet loss allows us to learn a ranking

$$\|f(A) - f(P)\|^2 < \|f(A) - f(N)\|^2$$

$$\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 < 0$$

$$\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + m < 0$$



margin

Schroff et al „FaceNet: a unified embedding for face recognition and clustering“. CVPR 2015

Triplet loss

- Triplet loss allows us to learn a ranking

$$\|f(A) - f(P)\|^2 < \|f(A) - f(N)\|^2$$

$$\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 < 0$$

$$\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + m < 0$$

$$\mathcal{L}(A, P, N) = \max(0, \|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + m)$$

Schroff et al „FaceNet: a unified embedding for face recognition and clustering“. CVPR 2015

Triplet loss

- Triplet loss allows us to learn a ranking
 - If $d(A, P) > d(A, N)$ the loss is going to be positive
 - If $d(A, P) < d(A, N)$ then I look at the margin

$$\mathcal{L}(A, P, N) = \max(0, ||f(A) - f(P)||^2 - ||f(A) - f(N)||^2 + m)$$

Schroff et al „FaceNet: a unified embedding for face recognition and clustering“. CVPR 2015

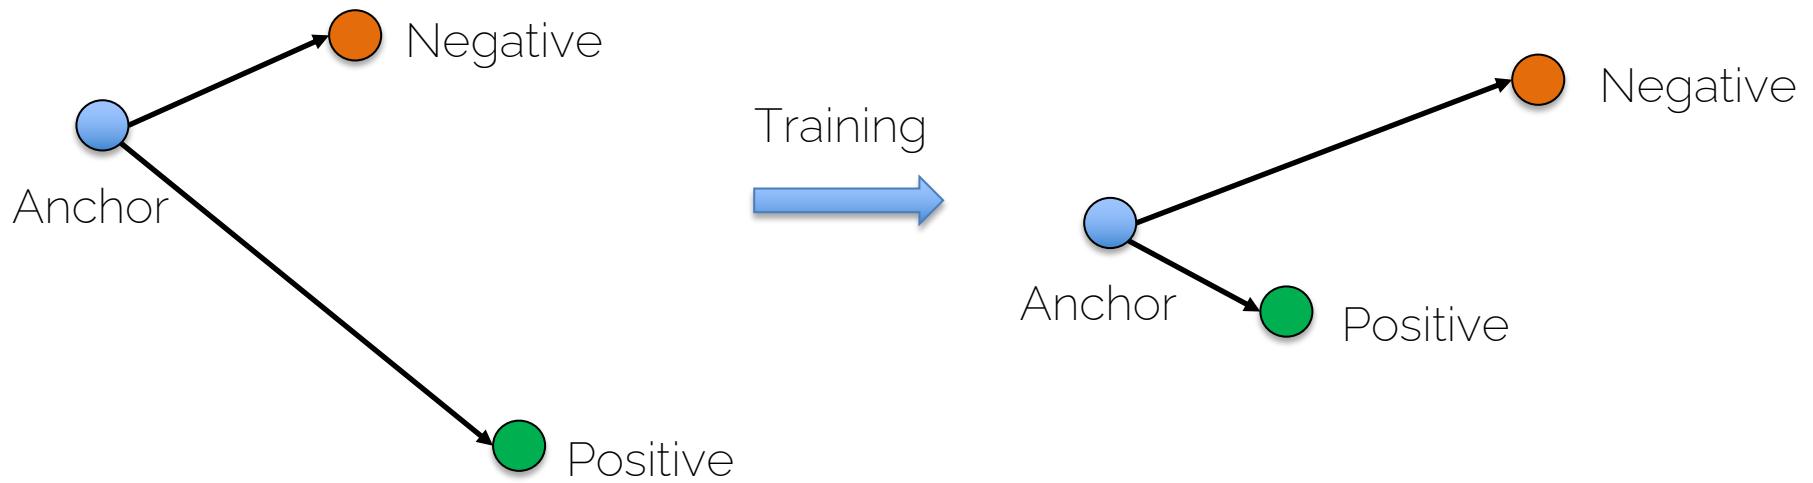
Triplet loss

- Training with hard cases

$$\mathcal{L}(A, P, N) = \max(0, \|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + m)$$

- Train for a few epochs
- Choose the hard cases where $d(A, P) \approx d(A, N)$
- Train with those to refine the distance learned

Triplet loss



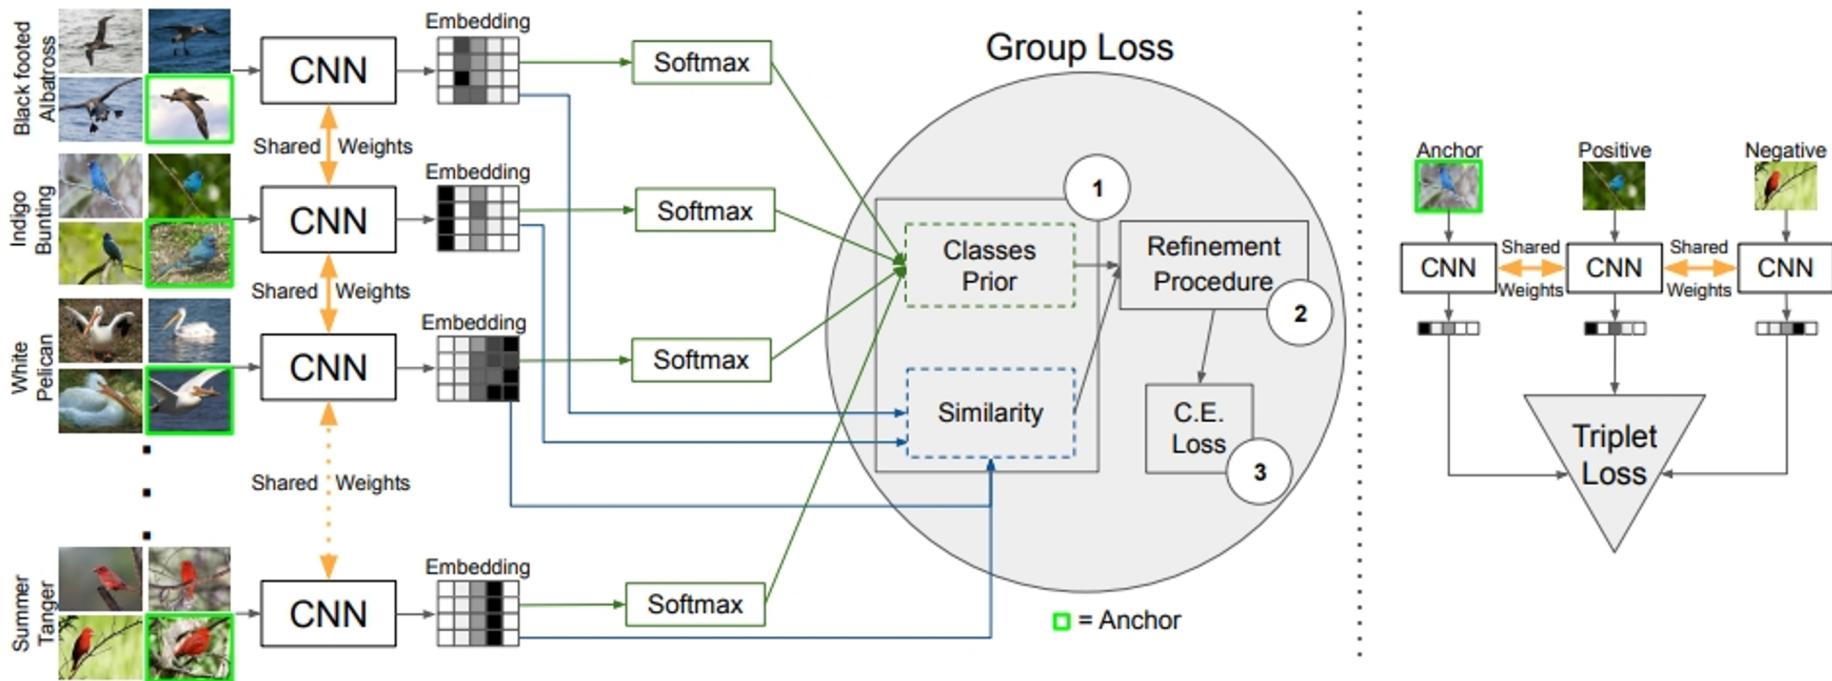
Problems with these losses

- Number of pairwise relations in a mini-batch is $O(n^2)$, but contrastive loss considers only $O(n/2)$ relations, and triplet loss considers only $O(2n/3)$ relations.
- Too much information is thrown away.
- In order to train these networks many tricks are required: hard-negative mining, intelligent sampling and multi-task learning.

Group Loss

Overview

We want to take into account the similarity of all samples wrt all other samples!

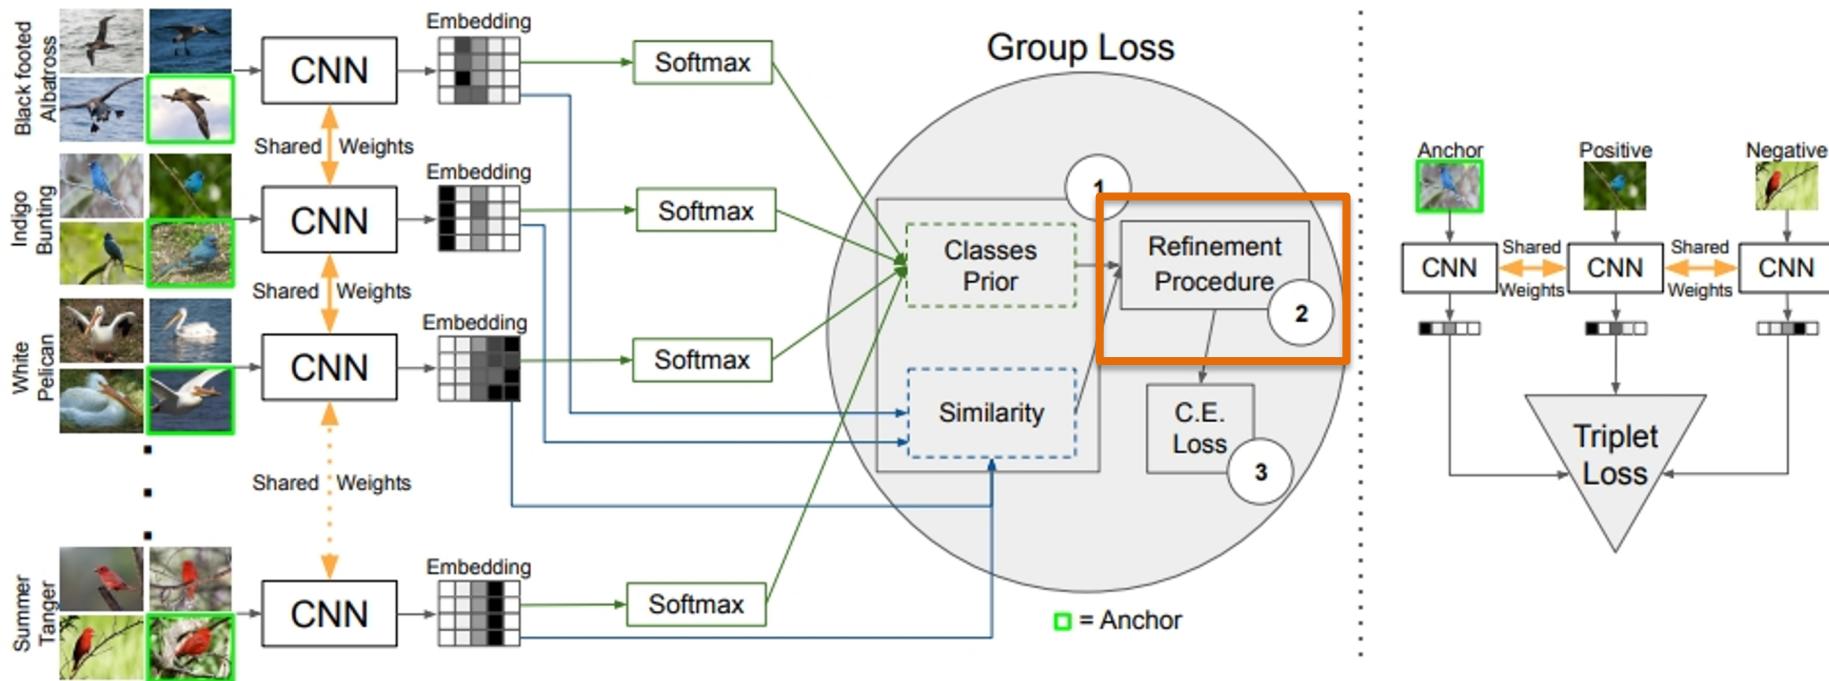


Elezi et al „The Group Loss for deep metric learning“. Arxiv: 1912.00385. 2019

Summary of the model

- 1) **Initialization:** Initialize X , the image-label assignment using the softmax outputs of the neural network. Compute the $n \times n$ pairwise similarity matrix W using the neural network embedding.
- 2) **Refinement:** Iteratively, refine X considering the similarities between all the mini-batch images, as encoded in W , as well as their labeling preferences.
- 3) **Loss computation:** Compute the cross-entropy loss of the refined probabilities and update the weights of the neural network using backpropagation.

Step 2: Refinement



Step 2: Refinement

B and C are much more similar than A and B, or A and C

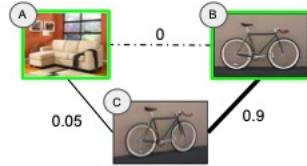
Consider a mini-batch with:

3 images (A,B,C)

4 classes (1,2,3,4)

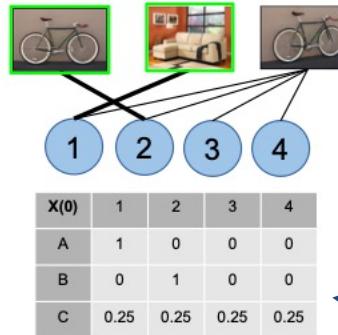
■ = Anchor

Batch Similarities



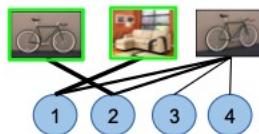
	W	A	B	C
A	0	0	0.05	
B	0	0	0.9	
C	0.05	0.9	0	

Labeling Priors



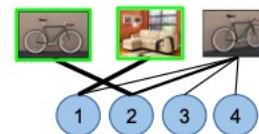
	1	2	3	4
X(0)	1	2	3	4
A	1	0	0	0
B	0	1	0	0
C	0.25	0.25	0.25	0.25

1st iteration



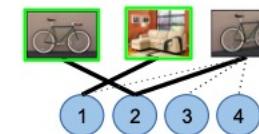
X(1)	1	2	3	4
A	1	0	0	0
B	0	1	0	0
C	0.28	0.39	0.16	0.17

2nd iteration



X(2)	1	2	3	4
A	1	0	0	0
B	0	1	0	0
C	0.20	0.67	0.06	0.07

3rd iteration



X(3)	1	2	3	4
A	1	0	0	0
B	0	1	0	0
C	0.03	0.95	0.01	0.01

Maybe the net is untrained and initialized the probabilities to an uniform distribution.

Step 2: Refinement - math

- Use replicator dynamics*

Iterate $t + 1$ times

Lambda is the class

$$x_{i\lambda}(t+1) = \frac{x_{i\lambda}(t)\pi_{i\lambda}(t)}{\sum_{\mu=1}^m x_{i\mu}(t)\pi_{i\mu}(t)}$$

Propagate information

Normalize in order to stay in standard simplex.

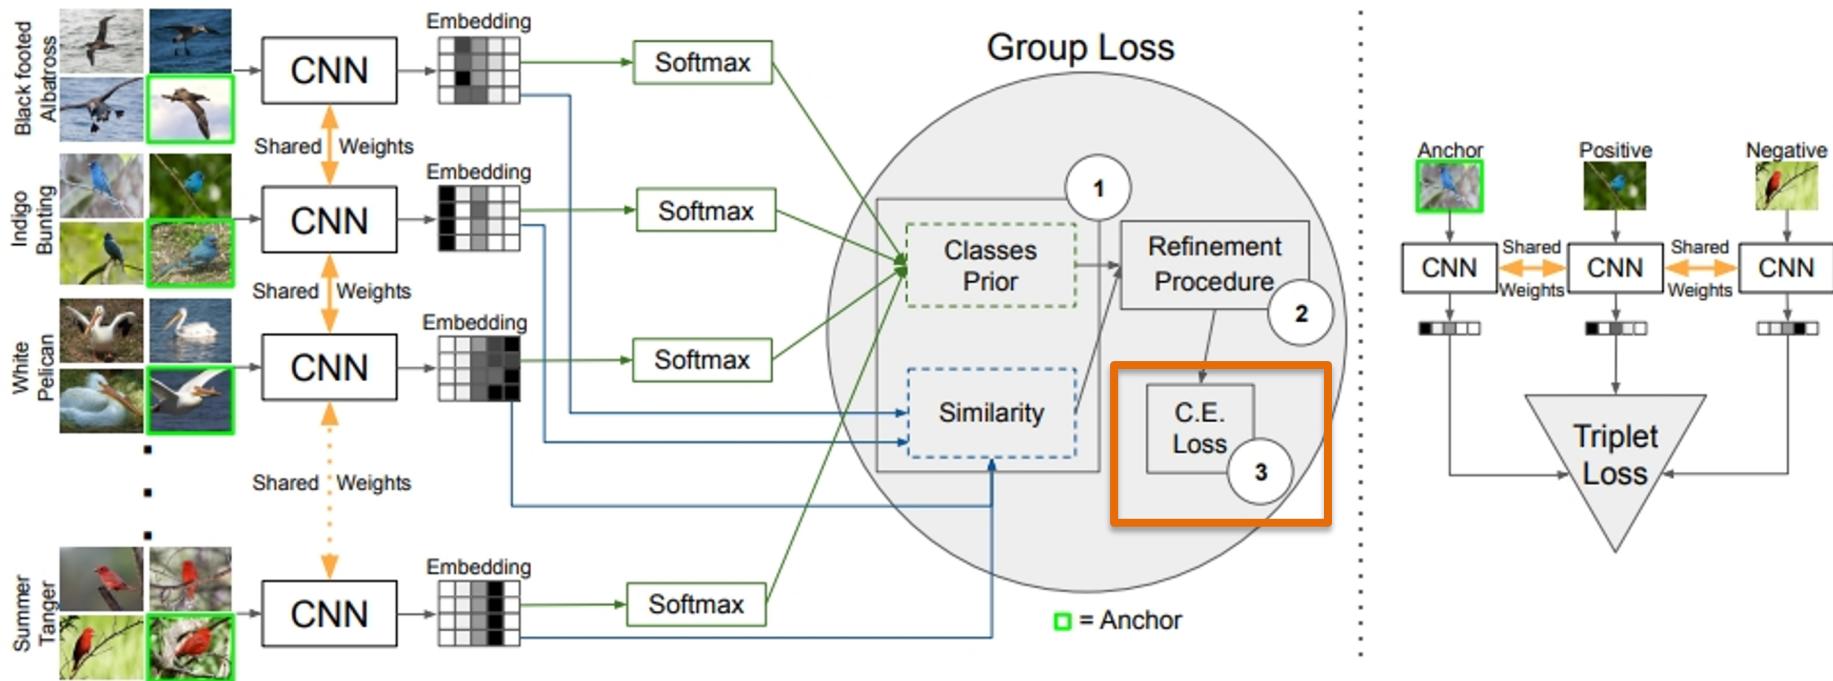
From the similarity matrix

$$\pi_{i\lambda} = \sum_{j=1}^n w_{ij} x_{j\lambda}$$

This measures the support that the current mini-batch gives to image I belonging to class lambda

*J.W. Weibull. Evolutionary Game Theory. MIT Press, 1997

Step 3: Loss and backprop



Step 3: Loss

- Compute cross-entropy over the refined probabilities.
- Backpropagate over the entire net
 - Group Loss has no parameters to learn, but it propagates the gradients over the network.
- Result: state-of-the-art results in retrieval

Elezi et al „The Group Loss for deep metric learning“. Arxiv: 1912.00385. 2019

Overview

- Mostly focused on online tracking
- Better appearance models with similarity learning (re-ID)
- How about offline tracking? Taking into account more than two frames to compute the matching (to recover from errors, missing detections, wrong predictions)

Lecture 5: MOT with graphs
- Better motion models?

Lecture 6: Trajectory prediction

Object tracking