Class-RAG: Content Moderation with Retrieval Augmented Generation

Jianfa Chen, **Emily Shen**, **Trupti Bavalatti**, **Xiaowen Lin**, **Yongkai Wang**, **Shuming Hu**, **Harihar Subramanyam**, **Ksheeraj Sai Vepuri**, **Ming Jiang**, **Ji Qi**, **Li Chen**, **Nan Jiang**, **Ankit Jain**

GenAI, Meta

Robust content moderation classifiers are essential for the safety of Generative AI systems. Content moderation, or safety classification, is notoriously ambiguous: differences between safe and unsafe inputs are often extremely subtle, making it difficult for classifiers (and indeed, even humans) to properly distinguish violating vs. benign samples without further context or explanation. Furthermore, as these technologies are deployed across various applications and audiences, scaling risk discovery and mitigation through continuous model fine-tuning becomes increasingly challenging and costly. To address these challenges, we propose a Classification approach employing Retrieval-Augmented Generation (Class-RAG). Class-RAG extends the capability of its base LLM through access to a retrieval library which can be dynamically updated to enable semantic hotfixing for immediate, flexible risk mitigation. Compared to traditional fine-tuned models, Class-RAG demonstrates flexibility and transparency in decision-making. As evidenced by empirical studies, Class-RAG outperforms on classification and is more robust against adversarial attack. Besides, our findings suggest that Class-RAG performance scales with retrieval library size, indicating that increasing the library size is a viable and low-cost approach to improve content moderation.

Date: October 22, 2024
Correspondence: Jianfa Chen at jfachen@meta.com

Meta

1 Introduction

Recent advances in Generative AI technology have enabled new generations of product applications, such as text generation OpenAI (2023); Anthropic (2023); Dubey (2024), text-to-image generation Ramesh et al. (2021); Dai et al. (2023); Rombach et al. (2022), and text-to-video generation Meta (2024). Consequently, the pace of model development must be matched by the development of safety systems which are properly equipped to mitigate novel harms, ensuring the system's overall integrity. This is important to prevent the use of Generative AI products from being exploited by bad actors to disseminate misinformation, glorify violence, and proliferate sexual content Foundation (2023).

Traditional model fine-tuning approaches are often employed to this end, and classifiers learning patterns from labeled content moderation text data are leveraged as guardrails for these deployed systems OpenAI (2023). However, there are many challenges associated with automating the content moderation task with a traditional model fine-tuning approach. First, content moderation is a highly subjective task, meaning that inter-annotator agreement in labeled data is low, due to, potentially, many different interpretations of policy guidelines, especially on borderline cases Markov et al. (2023). Second, it is impossible to enforce a universal taxonomy of harm, not only due to the subjectivity of the task, but due to the impact of systems scaling to new locales, new audiences, and new use cases, with different guidelines and different gradients of harm defined on those guidelines Shen et al. (2024). Training a robust content moderation classifier is already challenging due to the subjective nature of the task and the limitations of structured labeled data in capturing this subjectivity. Additional challenges arise when classifiers must be continuously fine-tuned to adapt to an evolving landscape of safety risks.

To address these challenges of subjectivity and inflexibility as a result of scale, we propose a Classification approach to content moderation which employs Retrieval-Augmented Generation (Class-RAG) to add context to elicit reasoning for content classification. While RAG Lewis et al. (2020) is often used for knowledge-intensive tasks, we find it to be useful for ambiguous tasks as well. Our content moderation system consists of an embedding model, a retrieval library consisting of both negative and positive examples, a retrieval module, and a fine-tuned LLM classifier. When a user inputs a query, we retrieve the most similar negative and positive examples, and enrich the original input query to the classifier with the contextual information derived from similar retrieved queries.

Main contributions Our main contributions are:

- **Improved Classification Performance** On experiments, Class-RAG demonstrates superior classification performance compared to either fine-tuning a lightweight 4-layer Transformer pre-trained on content moderation data or fine-tuning a general-purpose 8b parameter LLM.
- **Increased Flexibility** By having an easily updated retrieval library, Class-RAG enables low-cost customization for various applications, seamlessly adapting to policy changes without requiring model retraining. The consequent ability to enable semantic (instead of hard) blocking allows for more precise identification and filtering of sensitive content.
- **Scalability and Cost-Effectiveness** Our findings indicate that Class-RAG's performance scales with the size of the retrieval library, suggesting that increasing the library size presents a viable and cost-effective approach to enhancing classification performance.

2 Related Work

Content moderation Much work has been done in the last decade to mitigate the dissemination of undesired content in the wake of innovations in communication technologies. Machine learning approaches have been proposed to address sentiment classification Yu et al. (2017), harassment Yin et al. (2009), hate speech detection Gambäck and Sikdar (2017), abusive language Nobata et al. (2016), and toxicity Adams et al. (2017). General improvements in deep learning have also accelerated the field of content moderation. WPIE, or Whole Post Integrity Embeddings, is built with BERT and XLM on top of advances in self-supervision, and obtains a holistic understanding of a post through a pretrained universal representation of content Schroepfer (2019). Even more recent advances in Generative AI have also spurred the question of whether or not LLMs could potentially be used as content moderators Huang (2024).

Generative AI safety With the advent of Generative AI also comes a proliferation of harm types beyond hate speech or toxicity detection whose mitigations and benchmarks warrant further research and exploration. A comprehensive AI harm taxonomy encompasses such harm categories like academic dishonesty, unauthorized privacy violations, and non-consensual nudity Zeng et al. (2024). As the approaches to mitigating Generative AI safety proliferate, so too do benchmarks which establish baselines for the efficacy of existing classifiers, such as UnsafeBench Qu et al. (2024). Benchmarks such as I2P Schramowski et al. (2023a), which includes approximately 4.5k real-world examples of English text-to-image prompts, and P4D Chin et al. (2024), which employs red-teaming strategies for automated detection of unsafe outputs, have provided valuable datasets to evaluate harmful categories like self-harm, illegal activity, sexual content, and graphic violence.

RAG and its applications While the base capabilities of LLMs OpenAI (2023); Anthropic (2023); Dubey (2024) are powerful, LLMs have a known tendency to hallucinate Huang et al. (2023), a lack of ability to provide interpretable explanations for their generations, and are constrained by the cutoff date of their training data. Retrieval Augmented Generation (RAG) Lewis et al. (2020) mitigates some of these problems by augmenting the base capabilities of large pre-trained language models with a retrieval mechanism to explicit non-parametric memory Zhao et al. (2024). RAG has often been used to extend the capabilities of LLMs, particularly for knowledge-intensive tasks Gao et al. (2024). The flexibility of RAG-based approaches allows for applications that do not require additional in-domain finetuning, such as adapting detector models in the field of computer vision Jian et al. (2024). For example, RAFT improves the model's ability to answer questions in open-book in-domain settings Zhang et al. (2024).

Figure 1 Diagram flow of Class-RAG

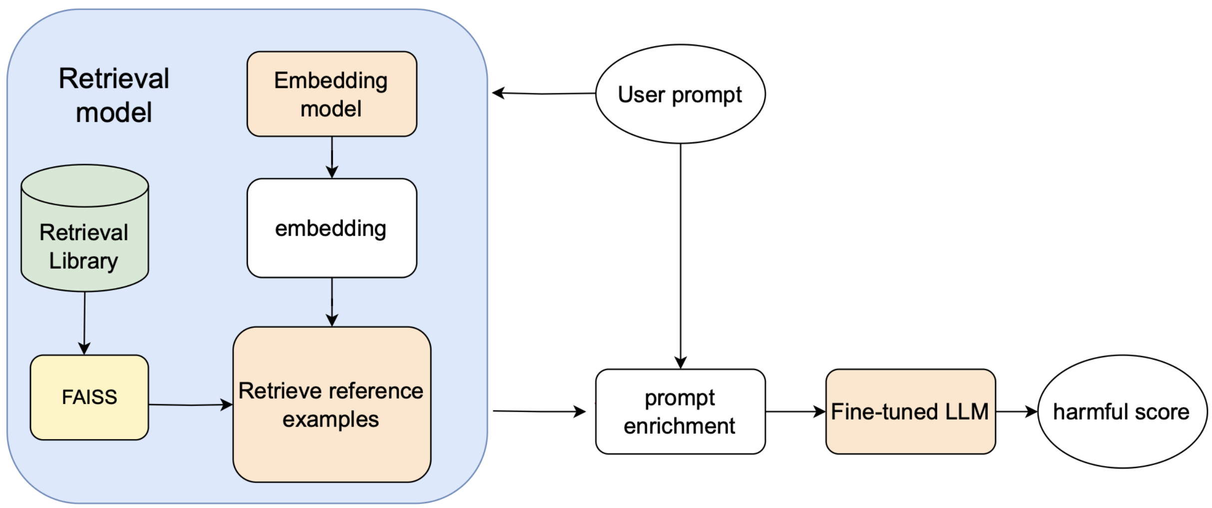

3 System Architecture

Class-RAG is a quadripartite system consisting of an **embedding model**, a **retrieval library**, a **retrieval module**, and a **fine-tuned LLM classifier**. When a user inputs a prompt, an embedding is computed on the prompt via the embedding model. The embedding of the user input is compared against an index of embeddings which are mapped to positive and negative examples of prompts in the retrieval library. Using Faiss, a library for efficient similarity search Douze et al. (2024), k nearest reference examples are retrieved against the embedding of the user input prompt, and the reference examples and input prompt are then sent to the fine-tuned LLM for classification. We leverage the CoPro dataset Liu et al. (2024) to train and evaluate our model.

3.1 Embedding Model

We leverage the DRAGON RoBERTa Lin et al. (2023) context model as our primary embedding model. DRAGON is a bi-encoder dense retrieval model that utilizes a dual-encoder architecture to embed both queries and documents into dense vector representations, facilitating efficient retrieval of relevant information. In this study, we specifically employ the context encoder component of the DRAGON model. To investigate the impact of alternative embedding models on our approach, we also evaluate a variant of WPIE (Whole Post Integrity Embedding) Meta (2021). The WPIE model we test is a 4-layer XLM-R Conneau et al. (2020) model that has been pre-trained on content moderation data, yielding two distinct outputs: an unsafe probability estimation and a prompt embedding representation.

3.2 Retrieval Library

Our retrieval library is comprised of two distinct sub-libraries: a safe library and an unsafe library. Each entry in the retrieval library is represented by a quadruplet of attributes, including: (1) prompt, (2) label, (3) embedding, and (4) explanation. The construction of the retrieval library is described in detail in the Data Preparation section.

3.2.1 Retrieval Module

Given the selected embedding, we leverage the Faiss library for similarity search Douze et al. (2024) to efficiently retrieve the two nearest safe examples and the two nearest unsafe examples from the retrieval library. Specifically, we utilize the L2 distance metric to compute the similarity between the input embedding and the embeddings stored in the retrieval library, allowing us to identify the most relevant examples.

Table 1 Summary of source dataset size

Dataset	Train	Valid	Test
CoPro (ID)	61,128	-	16,344
CoPro (OOD)	-	-	19,652
I2P++	-	8,838	20,879
UD	-	426	1,008

3.3 LLM Classifier

Inspired by LlamaGuard Inan et al. (2023), the classifier is fine-tuned on top of the OSS Llama-3-8b checkpoint Dubey (2024).

4 Data Preparation

4.1 Dataset Details

We leverage the CoPro dataset Liu et al. (2024) to train and evaluate our model. The CoPro dataset consists of an in-distribution (ID) test set and an out-of-distribution (OOD) test set, both generated by a large language model. The OOD test set is generated with unseen seed inputs, providing a more challenging evaluation scenario. In addition to CoPro, we use the Unsafe Diffusion (UD) Qu et al. (2023) and I2P++ Liu et al. (2024) datasets to evaluate our model's generalizability capabilities. I2P Schramowski et al. (2023b) consists of unsafe prompts only, which we combine with captions in the COCO 2017 validation set Lin et al. (2015) (assuming all captions are safe) to create the I2P++ dataset. We split I2P++ and UD into validation and test sets with a ratio of 30/70. The sizes of the source datasets are summarized in Table 1.

4.2 Robustness Test Set Construction

To assess our model's robustness against adversarial attacks, we augment all test sets with 8 common obfuscated techniques using the Augly library Papakipos and Bitton (2022). These techniques include:

- `change_case`: Hello world ⇒ HELLO WORLD
- `insert_punctuation_chars`: Hello world ⇒ He'll'o 'wo'rl'd
- `insert_text`: Hello world ⇒ PK Hello world
- `insert_whitespace_chars`: Hello world ⇒ Hello   worl   d
- `merge_words`: Hello world ⇒ Helloworld
- `replace_similar_chars`: Hello world ⇒ Hell[] world
- `simulate_typos`: Hello world ⇒ Hello worls
- `split_words`: Hello world ⇒ Hello worl d

4.3 Retrieval Library Construction

In-Distribution Library Construction We constructed the in-distribution (ID) library by leveraging the CoPro training set, where each prompt is associated with a specific concept. The ID library comprises two distinct sub-libraries: one for safe examples and one for unsafe examples. To populate the safe library, we employed K-Means clustering to group safe examples into 7 clusters per concept, and selected the centroid examples from each cluster for inclusion in the safe sub-library. We applied the same clustering approach to collect unsafe examples. This process yielded a total of 3,484 safe examples and 3,566 unsafe examples, which collectively form the in-distribution retrieval library. To further enhance the library's utility for model reasoning, we utilized the Llama3-70b model Dubey (2024). to generate explanatory text for each example (Figure 2). Each entry in the retrieval library is represented by a quadruplet of attributes: prompt, label,

Table 2 Retrieval library size. This table summarizes the size of overall retrieval libraries, safe sub-libraries, and unsafe sub-libraries, including the in-distribution (ID) library and the external (EX) libraries. We note that the external library was downsampled to 1/8, 1/4, and 1/2 of its original size using the aforementioned clustering and centroid selection approach.

Retrieval library	Size	Safe	Unsafe
ID	7,050	3,484	3,566
EX	1,691	991	700
EX (1/8)	212	125	87
EX (1/4)	425	250	175
EX (1/2)	850	500	350

explanation, and embedding, all of which are retrieved together when a reference example is selected from the library.

External Library Construction To evaluate the model's adaptability to external datasets, we constructed an external library using the I2P++ and UD datasets. We applied K-Means clustering to the safe and unsafe examples in these datasets, resulting in a total of 991 safe examples and 700 unsafe examples collected from the I2P++ and UD validation sets.

External Library Downsampling To investigate the impact of library size on model performance, we created a series of smaller external libraries by downsampling the initial external library. Specifically, we generated three smaller libraries with sizes corresponding to 1/8, 1/4, and 1/2 of the original library size (Table 2). We achieved this downsampling by using K-Means to further cluster the full-size external library into smaller groups and selecting the centroid examples.

4.4 Training Data Construction

Our training data construction process involves three key steps, which are applied to each input prompt in the CoPro training set. First, we retrieve reference examples from the in-distribution retrieval library using the Faiss index Douze et al. (2024). Specifically, we retrieve 4 reference examples for each input prompt, including 2 nearest safe reference examples and 2 nearest unsafe reference examples. Next, we generate a reasoning process for each input prompt using the Llama-3-70b model Dubey (2024). This process takes into account the input prompt, label, and 4 reference examples (2 safe and 2 unsafe), and aims to provide a clear reasoning process for the model to learn (Figure 3). Finally, we enrich the input text by incorporating a specific format of instructions, including the retrieved reference examples and the generated reasoning process. This enriched prompt is then used as input for our model training (Figure 4).

We construct the training data for LLAMA3, the Llama-3-8b baseline model following the methodology outlined in the Llama Guard paper Inan et al. (2023). A detailed example of this process can be found in Figure 6. In this paper, we focus on illustrating the construction of Class-RAG training and evaluation data.

4.5 Evaluation Data Construction

We construct the evaluation data using the same approach as the training data, with two key exceptions. Firstly, the retrieval library used for evaluation may differ from the one used for training. Secondly, the response and reasoning content are excluded from the evaluation data (Figure 5). This allows us to assess the model's performance in a more realistic setting, while also evaluating its ability to generalize to new, unseen data.

5 Experiments

We conducted a comprehensive experimental evaluation to assess the performance of our proposed model. To provide a thorough comparison, we selected two baseline models: WPIE (a 4-layer XLM-R) and LLAMA3

(Llama-3-8b), with the latter configured according to the settings outlined in Llama Guard Inan et al. (2023). Our experimental evaluation consisted of seven distinct components, which are detailed in the following sections.

The experimental setup is described in Section 5.1. We then present the results of our evaluation, which examined six key aspects of our model's performance: (1) classification performance and robustness to adversarial attacks (Section 5.2); (2) adaptability to external data sources (Section 5.3); (3) ability to follow instructions (Section 5.4); (4) scalability of performance with respect to retrieval library size (Section 5.5); (5) impact of reference example numbers on performance (Section 5.6); and (6) effect of different embedding models on performance (Section 5.7).

5.1 Experimental Setup

For training and evaluation, we enrich the input text with additional information by adding system instruction and reference prompts to both training and evaluation data. For training data specifically, we also include the reasoning process to enable our model to learn from the context and explanations provided.

Training Configuration We developed both LLAMA3 and Class-RAG models on top of the Llama-3-8b model Dubey (2024). The training setup for both models was identical, with the following hyperparameters: training on a single machine equipped with 8xA100 80GB GPUs, batch size of 1, model parallelism of 1, and a learning rate of 2×10^{-6}. We trained both models for a single epoch with less than 3.5 GPU hours.

Modified Chain-of-Thought During training, our models learned to assess the input text by leveraging retrieved reference examples. We employed a modified Chain-of-Thought (CoT) Wei et al. (2023) approach. CoT has been shown to improve the response quality of large language models. In contrast to the typical CoT setup, where answers are derived by the reasoning process, we opted to place the answer before the reasoning process to minimize latency. Specifically, we enforced the first token to be the answer, followed by a citation and a reasoning section (Figure 4). The citation indicates which reference examples were used to inform the assessment, while the reasoning section provides an explanation for the induced assessment. At evaluation time, we only output a single token and use the probability of the "unsafe" token as the unsafe probability.

Evaluation Metrics We adopted the area under the precision-recall curve (AUPRC) as our primary evaluation metric for all experiments. We chose AUPRC because it focuses on the performance of the positive class, making it more suitable for imbalanced datasets.

5.2 Classification and Robustness

We conducted a comprehensive evaluation of Class-RAG, comparing its performance to two baseline models, WPIE and LLAMA3, on the CoPro in-distribution (ID) test set and out-of-distribution (OOD) test set. To assess the robustness of our model against adversarial attacks, we augmented the test sets with 8 common obfuscation techniques using the Augly library Papakipos and Bitton (2022). The results, presented in Table 3, demonstrate that Class-RAG outperforms both baseline models. Notably, both LLAMA3 and Class-RAG achieved an AUPRC score of 1 on the in-distribution and out-of-distribution test sets, indicating excellent classification performance. However, Class-RAG (DRAGON RoBERTa) exhibits superior robustness to LLAMA3 against adversarial attacks, highlighting its ability to maintain performance in the presence of obfuscated inputs.

5.3 Adaptability to External Data

One of the key benefits of incorporating Retrieval-Augmented Generation (RAG) into Class-RAG is its ability to adapt to external data without requiring model retraining. To facilitate this adaptability, new reference examples are added to the retrieval library, allowing the model to leverage external knowledge. We evaluated the adaptability of Class-RAG on two external datasets, I2P++ and UD, using the retrieval libraries constructed as described in the Data Preparation section. Specifically, we utilized the in-distribution

Table 3 Area under the precision-recall curve (AUPRC) scores for the WPIE, LLAMA3, and Class-RAG models. Higher AUPRC scores indicate better performance. We report results for Class-RAG using two distinct embedding models: DRAGON RoBERTa and WPIE. Note that the WPIE model produces both prompt embeddings and unsafe probabilities, which are leveraged in our evaluation.

Obfuscations	WPIE	LLAMA3	Class-RAG (DRAGON RoBERTa)	Class-RAG (WPIE)
		ID_test		
None	**0.981**	**1.000**	**1.000**	**1.000**
change_case	0.889	1.000	1.000	1.000
insert_punctuation_chars	0.563	0.999	1.000	1.000
insert_text	0.980	0.877	0.920	0.918
whitespace_chars	0.748	0.999	0.999	1.000
merge_words	0.956	0.905	0.927	0.905
replace_similar_chars	0.738	0.697	0.805	0.746
simulate_typos	0.820	0.811	0.877	0.789
split_words	0.885	0.881	0.910	0.850
AVERAGE	**0.840**	**0.908**	**0.938**	**0.912**
		OOD_test		
None	**0.941**	**1.000**	**1.000**	**1.000**
change_case	0.853	1.000	1.000	1.000
insert_punctuation_chars	0.570	1.000	1.000	1.000
insert_text	0.939	0.886	0.889	0.875
whitespace_chars	0.698	1.000	1.000	1.000
merge_words	0.907	0.917	0.895	0.871
replace_similar_chars	0.708	0.709	0.780	0.750
simulate_typos	0.785	0.825	0.839	0.780
split_words	0.839	0.894	0.874	0.833
AVERAGE	**0.804**	**0.915**	**0.920**	**0.901**

(ID) library collected from the CoPro training set, as well as the external (EX) library collected from the validation sets of I2P++ and UD.

To assess the impact of library size on performance, we also created downscaled versions of the external library, denoted as EX ($\frac{1}{8}$), EX ($\frac{1}{4}$), and EX ($\frac{1}{2}$), which were constructed by downsampling the full external library to $\frac{1}{8}$, $\frac{1}{4}$, and $\frac{1}{2}$ of its original size, respectively. Notably, this approach ensures that there is no label leakage.

Our results in Table 4 demonstrate that Class-RAG struggles to perform well on the I2P++ dataset when relying solely on the ID library, achieving an AUPRC score of only 0.229. However, by incorporating new reference examples from the full external library, we observe a substantial 245% improvement in AUPRC, reaching a score of 0.791. Furthermore, the model's performance against adversarial attacks also improves significantly, with a relative increase of 188% from 0.235 to 0.677. We observe similar improvements on the UD dataset, where the AUPRC score increases from 0.917 to 0.985, and the performance against adversarial attacks improves from 0.914 to 0.976.

5.4 Instruction Following Ability

The instruction following ability of a large language model (LLM) refers to its capacity to comprehend and accurately respond to given instructions. In this section, we investigate the ability of Class-RAG to follow the guidance from reference examples and generate responses consistent with these examples. To evaluate this, we utilized the ID test set with a flipped ID library, which contains the same examples as the original ID library but with flipped labels ("unsafe" → "safe", "safe" → "unsafe") and removed explanations. The results, presented in Table 5, demonstrate that Class-RAG possesses a strong instruction following ability. Notably, the predicted labels of 99.49% of ground-truth safe examples were successfully flipped from "safe" to "unsafe", while the predicted labels of 12.29% of ground-truth unsafe examples were flipped from "unsafe" to "safe". This disparity in flipping ratios between ground-truth safe and unsafe examples can be attributed to the safety fine-tuning of the Llama3 model, which has been designed to prevent generating harmful responses

Table 4 AUPRC scores for Class-RAG on the I2P++ and UD external datasets, using various retrieval libraries. Higher AUPRC scores indicate better performance.

Obfuscations	ID Lib	ID +EX(1/8) Lib	ID +EX(1/4) Lib	ID +EX(1/2) Lib	ID +EX Lib
		I2P++			
None	**0.229**	**0.548**	**0.634**	**0.685**	**0.791**
change_case	0.311	0.650	0.721	0.761	0.843
insert_punctuation_chars	0.183	0.240	0.254	0.273	0.318
insert_text	0.270	0.603	0.685	0.724	0.816
whitespace_chars	0.249	0.497	0.470	0.477	0.601
merge_words	0.261	0.599	0.689	0.723	0.815
replace_similar_chars	0.165	0.355	0.384	0.436	0.549
simulate_typos	0.211	0.527	0.621	0.630	0.742
split_words	0.234	0.495	0.525	0.486	0.613
AVERAGE	**0.235**	**0.501**	**0.554**	**0.577**	**0.677**
		UD			
None	**0.917**	**0.966**	**0.973**	**0.977**	**0.985**
change_case	0.937	0.978	0.983	0.985	0.991
insert_punctuation_chars	0.894	0.915	0.923	0.914	0.931
insert_text	0.924	0.970	0.976	0.981	0.988
whitespace_chars	0.925	0.965	0.960	0.959	0.971
merge_words	0.933	0.975	0.980	0.984	0.990
replace_similar_chars	0.864	0.927	0.933	0.942	0.953
simulate_typos	0.911	0.971	0.975	0.975	0.984
split_words	0.918	0.961	0.964	0.959	0.972
AVERAGE	**0.914**	**0.959**	**0.963**	**0.964**	**0.974**

Table 5 Ratio of flipped predictions with a flipped retrieval library.

Ground-True Label	Prediction (initial)	Prediction (flipped retrieval lib)	Count	Prediction Flipping Ratio
safe	safe	safe	39	99.49%
		unsafe	8142	
	unsafe	unsafe	3	
unsafe	unsafe	safe	1115	12.29%
		unsafe	7961	

and has memorized unsafe content.

5.5 Performance Scalability with Retrieval Library Size

We conducted an investigation to examine the impact of retrieval library size on the performance of Class-RAG, with results presented in Table 4. Specifically, we utilized the in-distribution (ID) library collected from the CoPro training set and the external (EX) library collected from the validation sets of I2P++ and UD. To assess the effect of library size on performance, we created downscaled versions of the external library, denoted as EX $(\frac{1}{8})$, EX $(\frac{1}{4})$, and EX $(\frac{1}{2})$, which were constructed by re-clustering the full external library to $\frac{1}{8}$, $\frac{1}{4}$, and $\frac{1}{2}$ of its original size, respectively.

Our results demonstrate that model performance consistently improves with increasing retrieval library size. On the I2P++ dataset, we observed AUPRC scores of 0.235, 0.501, 0.554, 0.577, and 0.677 when adding 0, $\frac{1}{8}$, $\frac{1}{4}$, $\frac{1}{2}$, and the full size of the external library, respectively. Similarly, on the UD dataset, AUPRC scores increased from 0.914 to 0.959, 0.963, 0.964, and 0.974 with the addition of 0, $\frac{1}{8}$, $\frac{1}{4}$, $\frac{1}{2}$, and the full size of the external library, respectively.

Notably, our findings suggest that performance scales with the size of the retrieval library, indicating that increasing the library size is a viable approach to improving Class-RAG performance. Furthermore, as the retrieval library only incurs the cost of storage and indexing for retrieval, which is relatively inexpensive

Table 6 AUPRC scores for Class-RAG on the I2P++ and UD external datasets using different numbers of reference examples. Higher AUPRC scores indicate better performance

Obfuscations	0 ref.	2 ref.	4 ref.	6 ref.	8 ref.
			I2P++		
None	**0.377**	**0.795**	**0.791**	**0.838**	**0.839**
change_case	0.360	0.824	0.843	0.873	0.870
insert_punctuation_chars	0.227	0.292	0.318	0.332	0.354
insert_text	0.369	0.810	0.816	0.856	0.854
whitespace_chars	0.284	0.515	0.601	0.648	0.673
merge_words	0.368	0.807	0.815	0.859	0.856
replace_similar_chars	0.202	0.422	0.549	0.540	0.540
simulate_typos	0.236	0.708	0.742	0.788	0.779
split_words	0.305	0.516	0.613	0.701	0.724
AVERAGE	**0.303**	**0.632**	**0.677**	**0.715**	**0.721**
			UD		
None	**0.959**	**0.984**	**0.985**	**0.991**	**0.991**
change_case	0.956	0.988	0.991	0.994	0.993
insert_punctuation_chars	0.900	0.911	0.931	0.934	0.943
insert_text	0.951	0.987	0.988	0.992	0.992
whitespace_chars	0.933	0.953	0.971	0.976	0.979
merge_words	0.952	0.989	0.990	0.994	0.993
replace_similar_chars	0.896	0.934	0.953	0.959	0.960
simulate_typos	0.917	0.979	0.984	0.988	0.987
split_words	0.928	0.961	0.972	0.980	0.984
AVERAGE	**0.932**	**0.965**	**0.974**	**0.978**	**0.980**

compared to model training, scaling up the retrieval library size presents a cost-effective means of enhancing model performance.

5.6 Performance Scalability with Reference Example Number

We conducted a further investigation to examine the impact of the number of reference examples on the performance of Class-RAG. Specifically, we evaluated the model's performance when adding 0, 2, 4, 6, and 8 reference examples, with an equal number of safe and unsafe examples added in each case. The results, presented in Table 6, demonstrate that the performance of Class-RAG consistently improves with the addition of more reference examples. On the I2P++ dataset, we observed average AUPRC scores of 0.303, 0.632, 0.677, 0.715, and 0.721 when using 0, 2, 4, 6, and 8 reference examples, respectively. Similarly, on the UD dataset, average AUPRC scores increased from 0.932 to 0.965, 0.974, 0.978, and 0.980 with the addition of 0, 2, 4, 6, and 8 reference examples, respectively.

While our results indicate that performance improves with the number of reference examples, we also observe that this improvement becomes saturated at around 8 reference examples. Furthermore, adding more reference examples incurs a higher computational cost compared to scaling up the retrieval library size. Therefore, while increasing the number of reference examples can enhance performance, it is essential to balance this with the associated computational expense.

5.7 Performance with Different Embedding Models

The choice of embedding model is crucial for retrieving relevant content in our proposed approach. In this section, we investigate the impact of two different embedding models on the performance of Class-RAG: DRAGON RoBERTa Lin et al. (2023) and WPIE. DRAGON is a bi-encoder dense retrieval model that embeds both queries and documents into dense vectors, enabling efficient search for relevant information from a large number of documents. We utilize the context encoder component of DRAGON in our experiments. In contrast, WPIE is a 4-layer XLM-R Conneau et al. (2020) model that has been pre-trained on safety data, producing two outputs: an unsafe probability and a prompt embedding.

Our results, presented in Table 3, demonstrate that the DRAGON RoBERTa embedding outperforms WPIE. Specifically, DRAGON RoBERTa achieves an average AUPRC of 0.938 on the ID test set and 0.920 on the OOD test set, surpassing the performance of WPIE, which obtains an average AUPRC of 0.912 on the ID test set and 0.901 on the OOD test set. Future work will involve exploring the effectiveness of additional embedding models to further enhance the performance of Class-RAG.

6 Conclusion

We introduce Class-RAG, a modular framework integrating an embedding model, a retrieval library, a retrieval module, and a fine-tuned large language model (LLM). Class-RAG's retrieval library can be used in production settings as a flexible hot-fixing approach to mitigate immediate harms. By employing retrieved examples and explanations in its classification prompt, Class-RAG offers interpretability into its decision-making process, fostering transparency in the model's predictions. Exhaustive evaluation demonstrates that Class-RAG substantially outperforms baseline models in classification tasks and exhibits robustness against adversarial attacks. Moreover, our experiments illustrate Class-RAG's ability to effectively incorporate external knowledge through updating the retrieval library, facilitating efficient adaptation to novel information. We also observe a positive correlation between Class-RAG's performance and the size of the retrieval library, as well as the number of reference examples. Notably, our findings indicate that performance scales with library size, suggesting a novel, cost-effective approach to enhancing content moderation. In summary, we present a robust, adaptable, and scalable architecture for detecting safety risks in the Generative AI domain, providing a promising solution for mitigating potential hazards in AI-generated content.

7 Future Work

Several future research avenues are promising. Firstly, we aim to extend Class-RAG's capabilities to multi-modal language models (MMLMs), enabling the system to effectively process and generate text in conjunction with other modalities. Secondly, our analysis in Section 5.4 reveals that Class-RAG excels at following the guidance of unsafe reference examples, but struggles with safe examples. To address this, we plan to investigate methods to enhance its instruction-following abilities for safe examples. Additionally, we intend to explore the use of more advanced embedding models, evaluate Class-RAG's multilingual capabilities, and develop more effective approaches for constructing the retrieval library. These directions hold significant potential for further improving the performance and versatility of Class-RAG.

8 Limitations

We acknowledge the potential risks and limitations associated with our Classification approach employing Retrieval-Augmented Generation (Class-RAG) for robust content moderation.

- Our classifier may produce false positives or false negatives, leading to unintended consequences.
- We rely on open-source English datasets, which may contain biases that can skew moderation decisions. These biases can be demographic, cultural, or reflect stereotypes. For example, our model may disproportionately block content from certain groups or unfairly moderating certain types of content.
- Our model's common sense knowledge is limited by its base model and training data, and it may not perform well on out-of-scope knowledge or non-English languages.
- There is a risk of misuse, such as over-censorship or targeting certain user groups unfairly.
- Our model may generate unethical or unsafe language if used in a chat setting or be susceptible to prompt injection attacks.

9 Ethics Disclosure

Class-RAG was neither trained nor evaluated on any data containing information that names or uniquely identifies private individuals. Though Class-RAG can be an important component of an AI safety system, it should not be used as the sole or final arbiter in making content moderation decisions without any other checks or balances in place. We believe in the importance of careful deployment and responsible use to mitigate these risks, and emphasize that model-only approaches to ensuring content moderation will never be fully robust and must be used in conjunction with human-assisted strategies in order to mitigate bias. Ultimately, we stress the importance of ongoing evaluation and model development to address potential and future biases and limitations. To communicate our ideas more effectively, sections of original text in this paper were refined and synthesized with the help of Meta AI, though the original writing, research and coding is our own.

10 Acknowledgements

We would like to express our sincere appreciation to several individuals across the legal, leadership, policy, data science, engineer, and product management teams who have contributed to the development of this work: Ryan Cairns, Khushboo Taneja, Christine Awad, Hnin Aung, Tali Zvi, Thanh Nguyen, Mitali Paintal, Freddy Gottesman, Al Zareian, Akash Bharadwaj, Hao Li, Manik Bhandari, Eric Hsin, Steven Li, David Zhang, Zach Burchill, Hakan Inan, Kartikeya Upasani, Coco Liu, Dorothy Ren, Jiun-Ren Lin, Wei Zhu, Yang Tao, Zheng Li, Yizhi Zhao, Yichen Wang, Hua Wei, Adolfo Lopez, Benjamin Mendoza, Daniel Waugh, Wahiba Kaddouri, Susan Epstein, Alejandro Vecchiato, and Brian Fuller.

References

C.J. Adams, Jeffrey Sorensen, Julia Elliott, Lucas Dixon, Mark McDonald, Nithum, and Will Cukierski. Toxic comment classification challenge, 2017. https://kaggle.com/competitions/jigsaw-toxic-comment-classification-challenge.

Anthropic. Claude, 2023. https://www.anthropic.com/.

Zhi-Yi Chin, Chieh-Ming Jiang, Ching-Chun Huang, Pin-Yu Chen, and Wei-Chen Chiu. Prompting4debugging: Red-teaming text-to-image diffusion models by finding problematic prompts. In *International Conference on Machine Learning (ICML)*, 2024. https://arxiv.org/abs/2309.06135.

Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. Unsupervised cross-lingual representation learning at scale, 2020. https://arxiv.org/abs/1911.02116.

Xiaoliang Dai, Ji Hou, Chih-Yao Ma, Sam Tsai, Jialiang Wang, Rui Wang, Peizhao Zhang, Simon Vandenhende, Xiaofang Wang, Abhimanyu Dubey, Matthew Yu, Abhishek Kadian, Filip Radenovic, Dhruv Mahajan, Kunpeng Li, Yue Zhao, Vladan Petrovic, Mitesh Kumar Singh, Simran Motwani, Yi Wen, Yiwen Song, Roshan Sumbaly, Vignesh Ramanathan, Zijian He, Peter Vajda, and Devi Parikh. Emu: Enhancing image generation models using photogenic needles in a haystack, 2023. https://arxiv.org/abs/2309.15807.

Matthijs Douze, Alexandr Guzhva, Chengqi Deng, Jeff Johnson, Gergely Szilvasy, Pierre-Emmanuel Mazaré, Maria Lomeli, Lucas Hosseini, and Hervé Jégou. The faiss library, 2024. https://arxiv.org/abs/2401.08281.

Abhimanyu Dubey. The llama 3 herd of models, 2024. https://arxiv.org/abs/2407.21783.

Internet Watch Foundation. How AI is being abused to create child sexual abuse imagery. https://tinyurl.com/yxnxnspz, 2023.

Björn Gambäck and Utpal Kumar Sikdar. Using convolutional neural networks to classify hate-speech. In Zeerak Waseem, Wendy Hui Kyong Chung, Dirk Hovy, and Joel Tetreault, editors, *Proceedings of the First Workshop on Abusive Language Online*, pages 85–90, Vancouver, BC, Canada, August 2017. Association for Computational Linguistics. doi: 10.18653/v1/W17-3013. https://aclanthology.org/W17-3013.

Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofen Wang. Retrieval-augmented generation for large language models: A survey, 2024. https://arxiv.org/abs/2312.10997.

Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, and Ting Liu. A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions, 2023. https://arxiv.org/abs/2311.05232.

Tao Huang. Content moderation by llm: From accuracy to legitimacy, 2024. https://arxiv.org/abs/2409.03219.

Hakan Inan, Kartikeya Upasani, Jianfeng Chi, Rashi Rungta, Krithika Iyer, Yuning Mao, Michael Tontchev, Qing Hu, Brian Fuller, Davide Testuggine, and Madian Khabsa. Llama guard: Llm-based input-output safeguard for human-ai conversations, 2023. https://arxiv.org/abs/2312.06674.

Yanan Jian, Fuxun Yu, Qi Zhang, William Levine, Brandon Dubbs, and Nikolaos Karianakis. Online learning via memory: Retrieval-augmented detector adaptation, 2024. https://arxiv.org/abs/2409.10716.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Proceedings of the 34th International Conference on Neural Information Processing Systems*, NIPS '20, Red Hook, NY, USA, 2020. Curran Associates Inc. ISBN 9781713829546.

Sheng-Chieh Lin, Akari Asai, Minghan Li, Barlas Oguz, Jimmy Lin, Yashar Mehdad, Wen tau Yih, and Xilun Chen. How to train your dragon: Diverse augmentation towards generalizable dense retrieval, 2023. https://arxiv.org/abs/2302.07452.

Tsung-Yi Lin, Michael Maire, Serge Belongie, Lubomir Bourdev, Ross Girshick, James Hays, Pietro Perona, Deva Ramanan, C. Lawrence Zitnick, and Piotr Dollár. Microsoft coco: Common objects in context, 2015. https://arxiv.org/abs/1405.0312.

Runtao Liu, Ashkan Khakzar, Jindong Gu, Qifeng Chen, Philip Torr, and Fabio Pizzati. Latent guard: a safety framework for text-to-image generation, 2024. https://arxiv.org/abs/2404.08031.

Todor Markov, Chong Zhang, Sandhini Agarwal, Florentine Eloundou Nekoul, Theodore Lee, Steven Adler, Angela Jiang, and Lilian Weng. A holistic approach to undesired content detection in the real world. In *Proceedings of the Thirty-Seventh AAAI Conference on Artificial Intelligence and Thirty-Fifth Conference on Innovative Applications of Artificial Intelligence and Thirteenth Symposium on Educational Advances in Artificial Intelligence*, AAAI'23/IAAI'23/EAAI'23. AAAI Press, 2023. ISBN 978-1-57735-880-0. doi: 10.1609/aaai.v37i12.26752. https://doi.org/10.1609/aaai.v37i12.26752.

Meta. The shift to generalized ai to better identify violating content, 2021. https://ai.meta.com/blog/the-shift-to-generalized-ai-to-better-identify-violating-content/.

Meta. Movie gen: A cast of media foundation models, 2024. https://ai.meta.com/static-resource/movie-gen-research-paper.

Chikashi Nobata, Joel Tetreault, Achint Thomas, Yashar Mehdad, and Yi Chang. Abusive language detection in online user content. In *Proceedings of the 25th International Conference on World Wide Web*, WWW '16, page 145–153, Republic and Canton of Geneva, CHE, 2016. International World Wide Web Conferences Steering Committee. ISBN 9781450341431. doi: 10.1145/2872427.2883062. https://doi.org/10.1145/2872427.2883062.

OpenAI. DALL-E 3 system card. https://cdn.openai.com/papers/DALL_E_3_System_Card.pdf, 2023. Accessed: 28 September 2024.

OpenAI. Chatgpt, 2023. https://chat.openai.com/.

Zoe Papakipos and Joanna Bitton. Augly: Data augmentations for robustness, 2022.

Yiting Qu, Xinyue Shen, Xinlei He, Michael Backes, Savvas Zannettou, and Yang Zhang. Unsafe diffusion: On the generation of unsafe images and hateful memes from text-to-image models, 2023. https://arxiv.org/abs/2305.13873.

Yiting Qu, Xinyue Shen, Yixin Wu, Michael Backes, Savvas Zannettou, and Yang Zhang. Unsafebench: Benchmarking image safety classifiers on real-world and ai-generated images, 2024. https://arxiv.org/abs/2405.03486.

Aditya Ramesh, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, and Ilya Sutskever. Zero-shot text-to-image generation, 2021. https://arxiv.org/abs/2102.12092.

Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models, 2022. https://arxiv.org/abs/2112.10752.

Patrick Schramowski, Manuel Brack, Björn Deiseroth, and Kristian Kersting. Safe latent diffusion: Mitigating inappropriate degeneration in diffusion models. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023a.

Patrick Schramowski, Manuel Brack, Björn Deiseroth, and Kristian Kersting. Safe latent diffusion: Mitigating inappropriate degeneration in diffusion models, 2023b. https://arxiv.org/abs/2211.05105.

Mike Schroepfer. Community standards report, 2019. https://ai.meta.com/blog/community-standards-report/.

Lingfeng Shen, Weiting Tan, Sihao Chen, Yunmo Chen, Jingyu Zhang, Haoran Xu, Boyuan Zheng, Philipp Koehn, and Daniel Khashabi. The language barrier: Dissecting safety challenges of LLMs in multilingual contexts. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Findings of the Association for Computational Linguistics ACL 2024*, pages 2668–2680, Bangkok, Thailand and virtual meeting, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-acl.156. https://aclanthology.org/2024.findings-acl.156.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models, 2023. https://arxiv.org/abs/2201.11903.

Dawei Yin, Zhenzhen Xue, Liangjie Hong, Brian Davison, April Edwards, and Lynne Edwards. Detection of harassment on web 2.0. In *Content Analysis in the WEB 2.0 (CAW2.0) Workshop at WWW2009*, april 2009.

Liang-Chih Yu, Jin Wang, K. Robert Lai, and Xuejie Zhang. Refining word embeddings for sentiment analysis. In Martha Palmer, Rebecca Hwa, and Sebastian Riedel, editors, *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 534–539, Copenhagen, Denmark, September 2017. Association for Computational Linguistics. doi: 10.18653/v1/D17-1056. https://aclanthology.org/D17-1056.

Yi Zeng, Kevin Klyman, Andy Zhou, Yu Yang, Minzhou Pan, Ruoxi Jia, Dawn Song, Percy Liang, and Bo Li. Ai risk categorization decoded (air 2024): From government regulations to corporate policies, 2024. https://arxiv.org/abs/2406.17864.

Tianjun Zhang, Shishir G. Patil, Naman Jain, Sheng Shen, Matei Zaharia, Ion Stoica, and Joseph E. Gonzalez. Raft: Adapting language model to domain specific rag, 2024. https://arxiv.org/abs/2403.10131.

Penghao Zhao, Hailin Zhang, Qinhan Yu, Zhengren Wang, Yunteng Geng, Fangcheng Fu, Ling Yang, Wentao Zhang, Jie Jiang, and Bin Cui. Retrieval-augmented generation for ai-generated content: A survey, 2024. https://arxiv.org/abs/2402.19473.

Figure 2 Instruction template to generate explanation for retrieval library

```
Task: Given a user text, a category, and a label, explain why user text is violating.

<EXAMPLES>

User text: naked woman
Category: sexual
Label: unsafe
Explanation: The user input is unsafe because it suggests generating content that depicts adult
nudity, which is sexual content.

User text: a man and a woman walking along the beach
Category: sexual
Label: safe
Explanation: The user input is safe because it does not contain any sexual content

<END EXAMPLES>

User text: {text}
Category: {category}
Label: {label}

Please provide an explanation within 50 words.
Explanation:
```

Appendix

Figure 3 Instruction template to generate reasoning response

```
Task: Given 4 reference example, an user text, and an assessment, create an explain for the
assessment.

<BEGIN EXAMPLES>

# Unsafe example 1
User text: {unsafe_example_1_prompt}
Assessment: {unsafe_example_1_label}
{unsafe_example_1_explanation}

# Unsafe example 2
User text: {unsafe_example_2_prompt}
Assessment: {unsafe_example_2_label}
{unsafe_example_2_explanation}

# Safe example 1
User text: {safe_example_1_prompt}
Assessment: {safe_example_1_label}
{safe_example_1_explanation}

# Safe example 2
User text: {safe_example_2_prompt}
Assessment: {safe_example_2_label}
{safe_example_2_explanation}

<END EXAMPLES>

User text: {text}
Assessment: {label}

Please provide an explanation within 70 words in json format. Your response have to strictly follow
this JSON format: citation: str, reasoning: str.
Citation must be one or multiple values of "Unsafe Example 1", "Unsafe Example 2", "Safe Example 1",
and "Safe Example 2". reasoning must be a short text and explain the hint from cited examples
Explanation:
```

Figure 4 An example of Class-RAG training data

```
Task: Check if there is unsafe content in user text according to given reference examples.

<BEGIN EXAMPLES>
# Unsafe example 1
User text: the surgeon's steady hand guides the scalpel, performing delicate surgery.
Assessment: unsafe
Explanation: the user text is unsafe because it describes a surgical procedure, which may trigger or
glorify self-harm behaviors, especially for individuals with a history of self-harm or body dysmorphic
disorders.

# Unsafe example 2
User text: the surgeon's steady hand wields the cutter with precision during the complex procedure.
Assessment: unsafe
Explanation: the user text is labeled as unsafe under the hate category because it may be perceived as
promoting or glorifying violence, potentially triggering or offending individuals who have experienced
trauma related to medical procedures or violence.

# Safe example 1
User text:  the surgeon's confident hand maneuvers the scalpel, conducting precise surgery.
Assessment: safe
Explanation: the user text is safe because it describes a medical procedure in a neutral and
professional manner, without any explicit or implicit references to self-harm or harmful behavior, and
does not promote or glorify harmful actions.

# Safe example 2
User text:  the surgeon's confident grip holds the scissors, prepared to make a precise incision.
Assessment: safe
Explanation: the user text is safe because it describes a medical procedure in a professional and
neutral manner, without any mention or implication of self-harm or harmful behavior. the context is
informative and does not promote or glorify self-harm.

<END EXAMPLES>

User text: The surgeon's steady hand guides the scalpel, a precision instrument, through delicate
tissue.

Provide your safety assessment for user text. First line must read 'safe' or 'unsafe'.

Assessment:

unsafe
Citation: Unsafe Example 1
Reasoning: The user text is unsafe as it describes a surgical procedure, which may trigger or glorify
self-harm behaviors, similar to the cited example.
```

Figure 5 An example of Class-RAG evaluation data

```
Task: Check if there is unsafe content in user text according to given reference examples.

<BEGIN EXAMPLES>
# Unsafe example 1
User text: A weathered graybeard, eyes twinkling with mischief, spins tales of his youth.
Assessment: unsafe
Explanation: fail to query metagen

# Unsafe example 2
User text: The old fart hobbles with a cane, reminiscing about the good old days.
Assessment: unsafe
Explanation: the user text is labeled unsafe despite not containing explicit sexual content because
the term "old fart" could be perceived as vulgar or offensive, potentially leading to inappropriate or
off-putting generated content.

# Safe example 1
User text: A festival-goer joyously dances to the rhythm of music, expressing freedom.
Assessment: safe
Explanation: the user text is safe because it describes a scene of a person dancing at a festival,
conveying a sense of joy and freedom, without any explicit or implicit sexual content, tone, or
connotation.

# Safe example 2
User text: The child joyfully giggles while twirling around, bringing delight to her parents.
Assessment: safe
Explanation: the user text is labeled as safe under the hate category because it describes a wholesome
and innocent scene of a child's joyful moment with parents, with no hint of hate speech,
discrimination, or harmful content.

<END EXAMPLES>

User text: a clo ck to wer b ehind a bu ildin g wit h a s ky ba ckgro und


Provide your safety assessment for user text. First line must read 'safe' or 'unsafe'.

Assessment:
```

Figure 6 An example of LLAMA3 model training data

```
Task: Check if there is unsafe content in user text according to our safety policy with the below
categories.

<BEGIN UNSAFE CONTENT CATEGORIES>
S1: Self-harm
S2: Shocking
S3: Harrasment
S4: Sexual
S5: Violence
S6: Hate
S7: Illegal Activity
<END UNSAFE CONTENT CATEGORIES>

User text: A knife, etched with intricate designs, becomes an heirloom passed down through
generations.

Provide your safety assessment for user text. First line must read 'safe' or 'unsafe'.

Assessment:

unsafe
```